



Graph-Augmented Hybrid Intelligence for Topology Anomaly Detection and Correction in Power Distribution Networks with Explainable Inference

Hanbo Wang¹✉

1. School of Information and Software Engineering, East China Jiaotong University, Nanchang 330013, China

Received ; accepted

E-mail: x270704@163.com.

© Higher Education Press 2026

Abstract

Power distribution networks (PDNs) increasingly rely on accurate topology models for power flow analysis, fault localization, and protection coordination. Topology data integrates heterogeneous sources—Supervisory Control and Data Acquisition (SCADA) telemetry, Common Information Model (CIM)/Scalable Vector Graphics (SVG) engineering models, and switch signals—where errors manifest as five anomaly categories: model-graph mismatch, topology interruption, virtual/faulty connections, measurement-topology inconsistency, and signal-measurement contradiction. Existing methods operate within a single paradigm and address only a subset of these categories. We propose a *three-layer hybrid intelligence framework* that unifies symbolic rule reasoning, physics-informed state estimation, and graph neural network (GNN) adaptive detection into an end-to-end pipeline of detection, localization, correction, and verification. A key contribution is the *differential physics detector*, which compares real-time telemetry against Newton–Raphson power flow expectations, providing a scale-invariant detection threshold. Cross-layer feature fusion enables the GNN to leverage state estimation residuals as supplementary input, achieving detection capability beyond any single layer. On a benchmark of 30 networks (3–1 888 buses, 5 independent seeds), the framework achieves 78.1% per-anomaly recall, a 2.4× improvement over Isolation Forest, and 99.4% network-level recall with near-zero variance ($\pm 0.5\%$). An improved multi-network GNN further increases recall to 83.5%, and extended evaluation on 39 networks with 28 anomaly subtypes confirms 96.1% overall recall. The macro-averaged precision of 44.8% reflects a deliberate recall-oriented design for safety-critical monitoring, where missed anomalies carry far higher costs than false alarms. Hard anomaly ablation demonstrates layer-specific contributions (state estimation (SE): –37 pp, GNN: –10 pp), confirming complementary rather than redundant coverage. The framework maintains 100% recall under noise up to 5%, completes inference in 0.94 s on average, and generates human-readable explanations via an integrated explainability module. Statistical significance is verified through paired Wilcoxon signed-rank tests ($p < 0.001$).

Key words

graph neural network; topology anomaly detection; hybrid intelligence; power distribution network; state estimation; multi-paradigm integration; explainable inference; data quality assurance

■ 1 Introduction

The transition toward carbon-neutral energy systems is accelerating the deployment of distributed energy resources (DERs)—including rooftop photovoltaics, battery storage, and electric vehicle charging stations—into power distribution networks (PDNs). By the end of 2025, China’s distributed photovoltaic capacity alone exceeded 300 GW [1]. This growth transforms PDNs from static radial structures into dynamic mesh-capable networks with frequent switching and bidirectional power flow [2, 3]. Consequently, topology models must integrate heterogeneous sources—SCADA telemetry, CIM/SVG engineering models, and switch signals—into a single coherent representation [4–6]. Each source carries different error modes and update frequencies, creating a critical data quality challenge for downstream applications such as power flow analysis and fault localization.

From a computer science perspective, this problem presents a unique challenge at the intersection of graph-based machine learning, symbolic reasoning, and physics-informed computation: reconciling heterogeneous data representations under physical constraints, where no single algorithmic paradigm provides complete coverage.

Accurate topology models underpin downstream applications including power flow analysis, state estimation, fault location, and protection coordination [7, 8]. In practice, PDN topology data flows through a complex pipeline involving SCADA (Supervisory Control and Data Acquisition) telemetry, CIM (Common Information Model) engineering models [1, 9], SVG (Scalable Vector Graphics) graphical representations, and manual maintenance by field engineers. Errors at each stage manifest as five topology anomaly categories:

1. **Model-Graph Mismatch (MGM):** CIM model devices do not

correspond one-to-one with SVG graphical elements;

2. **Topology Interruption (TI):** A disconnection exists between feeder-end devices and the power source;
3. **Virtual/Faulty Connections (VFC):** Device connectivity relationships do not match physical reality;
4. **Measurement-Topology Inconsistency (MTI):** SCADA telemetry data contradicts the topology model;
5. **Signal-Measurement Contradiction (SMC):** Switch status signals conflict with telemetry-inferred states.

Existing approaches for detecting and correcting these anomalies fall into three categories. *Rule-based methods* [7] employ graph-theoretic analysis and data consistency checks. They are fast and interpretable but cover only deterministic structural anomalies (MGM and TI). *Physics-based methods* [10–13] use weighted least squares (WLS) state estimation with statistical bad data detection. These are theoretically rigorous but require sufficient measurement redundancy and cannot directly identify graphical data errors. *Data-driven methods* [14–18] leverage graph neural networks (GNNs) to learn anomaly patterns. They detect complex composite anomalies but suffer from limited interpretability and poor generalization to unseen topologies.

No single existing method covers all five anomaly types while also providing automated correction and explainable results. The fundamental challenge is that each anomaly type requires a different detection paradigm: structural anomalies need symbolic reasoning, measurement anomalies need physics-based estimation, and pattern-based anomalies need data-driven learning. Existing hybrid approaches [19–21] combine at most two paradigms and do not address the full taxonomy. Recent hybrid AI advances [21–23] improve robustness by combining physics and data-driven methods. However, these efforts focus on power flow and state estimation rather than comprehensive topology anomaly detection and correction.

1.1 Contributions

We propose a *three-layer hybrid intelligence framework* that bridges symbolic reasoning, physics-informed computation, and deep learning for comprehensive PDN topology anomaly detection and correction. The framework makes both *algorithmic contributions* (differential physics detector, hierarchical detection, explainability module) and *empirical contributions* (hard anomaly benchmark, multi-seed validation, extended subtype analysis). We make the following contributions.

1. **Three-layer hybrid detection architecture.** We design a modular architecture integrating three detection paradigms at different time scales: (i) a rule engine for millisecond-level deterministic screening, (ii) a physics-informed state estimation layer with a differential physics detector, and (iii) a GNN-based adaptive detector. Cross-layer feature fusion enables the GNN to leverage state estimation residuals as supplementary input. The framework processes CIM/RDF models, SVG graphics, and SCADA telemetry through a unified detection–localization–correction–verification pipeline.

2. **Differential physics detector.** We introduce a detection principle that compares SCADA measurements against Newton–Raphson power flow expectations on the original topology. This provides a physics-grounded, scale-invariant threshold requiring no per-network calibration.
3. **Hierarchical detection for large-scale networks.** We propose a graph-coarsening-based partitioning strategy that combines connected component decomposition, voltage-level grouping, and Louvain community detection [24] to enable efficient detection on networks with up to 1 888 buses, reducing computational complexity from $O(n^3)$ to $O(k \cdot (n/k)^3)$ for k partitions.
4. **Explainability module.** We develop a perturbation-based feature importance analysis and reasoning chain explanation system that provides interpretable detection results without requiring external dependencies such as SHAP [25], addressing the critical interpretability requirements of power system operations.
5. **Closed-loop detect–correct–verify pipeline.** We implement an automated correction engine following the minimum-modification principle, with post-correction verification through power flow re-computation, enabling fully automated topology data quality assurance.

Section 2 reviews related work. Sections 3–4 present the problem formulation and framework architecture. Sections 5–6 describe hierarchical detection and explainability. Section 7 reports experiments, Section 8 discusses findings, and Section 9 concludes.

■ 2 Related Work

2.1 Graph Neural Networks for Power Systems

Graph neural networks have emerged as a powerful paradigm for power system analytics due to the inherent graph structure of electrical networks. Hamilton et al. [14] introduced GraphSAGE for inductive representation learning on large graphs, which has become a foundational architecture for graph-structured learning in power systems. The comprehensive survey by Wu et al. [16] categorizes GNN architectures and their applications across domains. Physics-informed machine learning approaches [26] have demonstrated that embedding physical constraints into neural network architectures improves generalization and interpretability in power system applications. A recent IEEE review [27] systematically examines GNN applications in state estimation, stability assessment, fault detection, renewable energy integration, dynamic topology management, and cyber-physical security, identifying topology-aware architectures as a key research direction. Recent surveys [28] have cataloged anomaly detection methods for power grids, highlighting the growing adoption of deep learning approaches.

GNNs have been applied to three main tasks in power systems. *Power flow and state estimation:* physics-informed GNNs embed Kirchhoff constraints into the loss function, achieving strong accuracy on IEEE benchmarks [17, 18, 29]. Meta-PIGAT [21] integrates meta-learning with physics-informed graph attention. Lopez et al. [30] demonstrated GNN-based AC power flow prediction. Wei et al. [31]

proposed joint topology and parameter identification via pseudo-label-trained GNNs. Dynamic SE approaches [32] have examined real-time state estimation under cyber-physical threats. *Topology identification*: multisource data fusion [33], edge-enhanced graph attention [19], and wavelet-based graph networks [34, 35] have achieved 85–92% accuracy on IEEE test systems. Knowledge-graph-augmented GNNs [36] improve fault tolerance in data-sparse scenarios. *Anomaly detection*: hybrid LSTM-GIN models [37], Bayesian GANs [38], and tensor decomposition [39] address false data injection and SCADA anomalies. Explainable GNN approaches [5, 40] have shown that interpretable reasoning chains improve operator trust. Recent surveys [28, 41] catalog the growing adoption of deep learning for grid analytics.

Despite these advances, existing GNN-based approaches address isolated aspects of the topology data quality problem rather than providing a unified framework. This work addresses this gap by integrating GNN-based detection as the third layer of a hybrid framework that also includes rule-based and physics-based layers.

2.2 Physics-Based State Estimation and Bad Data Detection

State estimation has underpinned power system monitoring since the seminal work of Schweppe et al. [10]. The weighted least squares (WLS) estimator [42] remains the standard approach, with the Gauss–Newton iterative method solving the normal equations:

$$\Delta \mathbf{x} = (\mathbf{H}^T \mathbf{R}^{-1} \mathbf{H})^{-1} \mathbf{H}^T \mathbf{R}^{-1} [\mathbf{z} - \mathbf{h}(\mathbf{x}^k)], \quad (1)$$

where $\mathbf{H}$ is the Jacobian matrix, $\mathbf{R}$ is the measurement covariance matrix, $\mathbf{z}$ is the measurement vector, and $\mathbf{h}(\cdot)$ is the nonlinear measurement function.

Comprehensive reviews [43] have cataloged anomaly detection and classification methods for state estimation, categorizing approaches from classical statistical tests to modern machine learning. Bad data detection (BDD) employs the χ^2 test [12] for global detection and normalized residual analysis [13] for localization:

$$J(\hat{\mathbf{x}}) = [\mathbf{z} - \mathbf{h}(\hat{\mathbf{x}})]^T \mathbf{R}^{-1} [\mathbf{z} - \mathbf{h}(\hat{\mathbf{x}})], \quad (2)$$

with $J(\hat{\mathbf{x}}) > \chi^2_{\alpha}(m - n)$ indicating the presence of bad data. Robust alternatives include LAV estimation [11], which minimizes $\sum |r_i|$ and naturally suppresses outliers, and IRLS with Huber weight functions [44]. Topology error identification uses hypothesis testing [12], enumerating candidate topologies and selecting the one minimizing $J(\hat{\mathbf{x}})$. Recent work [45] has developed comprehensive methods for topology error detection in state estimation using hypothesis testing enumeration, while robust estimators [46] based on nonlinear least absolute value have been proposed for simultaneous topology error identification and state estimation.

2.3 Topology Identification and Data Fusion

Topology identification in distribution networks has been extensively studied. Deka and Misra [3] provided a comprehensive tutorial on learning distribution grid topologies from voltage measurements. The multi-source data fusion approach by Wang et al. [47] reconstructed grid topology dynamically by combining heterogeneous data streams. Comprehensive overviews [48] have documented the evolution of

power system topology analysis from manual inspection to automated methods, identifying key challenges in scalability and robustness. Baldi et al. [2] surveyed AI methods for distribution network topology identification, categorizing approaches by input data type and identification objective.

Recent work has also explored physics-informed reinforcement learning for topology control [49] and differentiable power flow for topology optimization [23]. However, these approaches focus on *operational* topology control rather than *data quality* assurance, which is the primary concern of this paper.

2.4 Benchmarks and Datasets

Standard test feeders including IEEE 13/34/123/8500 bus systems [50, 51] and CIGRE MV/LV benchmarks [52] serve as the primary evaluation platforms. SimBench [53] provides 123 network models covering German MV/LV/EHV/HV voltage levels, while PandaPower [54] offers 43 built-in networks with integrated power flow and state estimation solvers. The PowerGraph benchmark [55] introduced a standardized GNN evaluation dataset for power systems.

2.5 Research Gap

Table 1 summarizes the positioning of existing approaches against the requirements of comprehensive topology data quality assurance. Pure rule-based methods [7] cover only deterministic structural anomalies. Physics-based state estimation [10, 11] addresses measurement anomalies but requires sufficient observability. Data-driven GNN methods [17, 18] show promise but lack built-in correction and explainability. Recent hybrid approaches [21, 22, 56] combine two paradigms but do not address the full anomaly taxonomy or provide closed-loop verification. No existing method simultaneously addresses four requirements of a production-grade topology data quality system: (i) full coverage of all five anomaly types, (ii) automated correction with closed-loop verification, (iii) explainable detection results for operator trust, and (iv) scalability to networks with thousands of buses. We address this gap through a modular three-layer architecture.

Table 1 Comparison of existing approaches against topology data quality requirements

Method	MGM	TI	VFC	MTI	SMC
Rule-based [7]	✓	✓	—	—	—
WLS-SE [10]	—	—	✓	✓	✓
Robust SE [11]	—	—	✓	✓	—
GNN-based [18]	—	—	✓	✓	—
Hybrid (2-layer) [56]	—	✓	✓	✓	—
Mechanism-Data [29]	—	—	✓	✓	—
Proposed (3-layer)	✓	✓	✓	✓	✓

3 Problem Formulation

3.1 Network Model

A power distribution network (PDN) is modeled as a graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$, where $\mathcal{V}$ denotes the set of buses (nodes) and $\mathcal{E}$ denotes the set of branches (edges) including lines, transformers, and switches. Each bus $v_i \in \mathcal{V}$ has an associated state vector $\mathbf{x}_i = [V_i, \theta_i]^T$ comprising

voltage magnitude and phase angle. The network model is described in CIM/Resource Description Framework (RDF) format [1, 57], while the graphical representation uses SVG with device-level annotations.

3.2 Measurement Model

SCADA measurements form a vector $\mathbf{z} \in \mathbb{R}^m$ follows the standard nonlinear model:

$$\mathbf{z} = \mathbf{h}(\mathbf{x}) + \mathbf{v}, \quad \mathbf{v} \sim \mathcal{N}(\mathbf{0}, \mathbf{R}), \quad (3)$$

where $\mathbf{h} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is the measurement function, $\mathbf{x} \in \mathbb{R}^n$ is the state vector with dimension $n = 2N$ for N buses (or $6N$ for three-phase models), and $\mathbf{v}$ is the measurement noise with covariance $\mathbf{R} = \text{diag}(\sigma_1^2, \dots, \sigma_m^2)$.

3.3 Anomaly Types

We formalize the five anomaly types as deviations from the ground-truth topology $\mathcal{G}^*$ and measurement vector $\mathbf{z}^*$:

Definition 1 (Model-Graph Mismatch). An MGM anomaly exists when the device sets $\mathcal{D}_{\text{CIM}}$ and $\mathcal{D}_{\text{SVG}}$ are inconsistent: $\mathcal{D}_{\text{CIM}} \neq \mathcal{D}_{\text{SVG}}$.

Definition 2 (Topology Interruption). A TI anomaly exists when $\mathcal{G}$ contains connected components C_k such that no element of C_k is connected to a power source bus.

Definition 3 (Virtual/Faulty Connection). A VFC anomaly exists when the connectivity relation $E \subset \mathcal{V} \times \mathcal{V}$ differs from the physical wiring E^* .

Definition 4 (Measurement-Topology Inconsistency). An MTI anomaly exists when measurements $\mathbf{z}$ are inconsistent with the electrical model implied by $\mathcal{G}$, i.e., $J(\hat{\mathbf{x}}) > \chi_\alpha^2(m - n)$ under correct topology.

Definition 5 (Signal-Measurement Contradiction). An SMC anomaly exists when the switch status vector $\mathbf{s}_{\text{signal}}$ from SCADA contradicts the status $\mathbf{s}_{\text{inferred}}$ inferred from power flow analysis.

3.4 Objective

Given the observed topology $\mathcal{G}$, measurement vector $\mathbf{z}$, and switch signal $\mathbf{s}_{\text{signal}}$, the objective is to: (i) detect all anomalies $A = \{a_1, \dots, a_K\}$; (ii) localize each anomaly to specific devices; (iii) generate a correction set $C = \{c_1, \dots, c_L\}$ minimizing $|C|$ subject to electrical constraints; and (iv) verify that applying C restores $\mathcal{G}$ to a valid state.

4 The Proposed Framework

4.1 Architecture Overview

Figure 1 illustrates the overall architecture of our three-layer hybrid intelligence framework. The system processes CIM/RDF models, SVG graphics, and SCADA telemetry through a data alignment module, then feeds the unified representation into three detection layers operating at different time scales and abstraction levels.

The cross-layer integration module merges detections using confidence-weighted fusion, and the correction engine generates minimal-modification proposals that are verified through closed-loop power flow re-computation.

Design rationale. The three-layer architecture embodies *defense in depth*: each layer operates independently with different data requirements. Layer 1 requires only structural data (graph topology, device

mappings). Layer 2 additionally requires electrical measurements. Layer 3 requires training data. This design ensures graceful degradation: if SCADA data is unavailable, Layer 1 still provides structural detection; if the GNN is untrained, Layers 1–2 still function. This redundant coverage provides operational robustness beyond single-layer approaches.

4.2 Layer 1: Rule Engine (Symbolic Reasoning)

The rule engine performs fast, deterministic screening of structural anomalies through graph-theoretic analysis and data consistency checks. It operates in $O(|V| + |E|)$ time and covers MGM and TI anomalies.

4.2.1 Model-Graph Mismatch Detection

MGM detection performs bidirectional set comparison between CIM model devices $\mathcal{D}_{\text{CIM}}$ and SVG graphical elements $\mathcal{D}_{\text{SVG}}$:

$$\text{MGM}_{\text{CIM-only}} = \{d \mid d \in \mathcal{D}_{\text{CIM}} \setminus \mathcal{D}_{\text{SVG}}\}, \quad (4)$$

$$\text{MGM}_{\text{SVG-only}} = \{d \mid d \in \mathcal{D}_{\text{SVG}} \setminus \mathcal{D}_{\text{CIM}}\}, \quad (5)$$

$$\text{MGM}_{\text{type}} = \{d \mid d \in \mathcal{D}_{\text{CIM}} \cap \mathcal{D}_{\text{SVG}}, T_{\text{CIM}}(d) \neq T_{\text{SVG}}(d)\}. \quad (6)$$

4.2.2 Topology Interruption Detection

TI detection identifies disconnected components lacking power source connectivity. Let $\mathcal{S} \subset \mathcal{V}$ denote the set of source buses. We compute connected components $\{C_1, \dots, C_K\}$ of $\mathcal{G}$ and flag component C_k as interrupted if $C_k \cap \mathcal{S} = \emptyset$.

4.2.3 Quick Physical Range Check

The rule engine also performs fast range-based validation of measurements without requiring state estimation: $V_i \in [0.93, 1.07]$ pu, $|P_{ij}| \leq P_{ij}^{\max}$, and switch signal consistency checks (multiple SCADA signals for the same switch must agree).

4.3 Layer 2: Physics-Informed State Estimation

Layer 2 employs physics-based state estimation for rigorous detection of measurement anomalies, faulty connections, and switch status inconsistencies. It operates in $O(n^3)$ per iteration (typically 3–5 iterations) and covers VFC, MTI, and SMC anomalies.

4.3.1 WLS State Estimation

The standard WLS estimator solves:

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x}} [\mathbf{z} - \mathbf{h}(\mathbf{x})]^T \mathbf{R}^{-1} [\mathbf{z} - \mathbf{h}(\mathbf{x})], \quad (7)$$

using the Gauss–Newton iterative method. We employ PandaPower’s [54] built-in state estimation module, which supports three-phase unbalanced modeling.

4.3.2 Bad Data Detection

Global detection uses the χ^2 test: if $J(\hat{\mathbf{x}}) > \chi_\alpha^2(m - n)$, bad data is present. Localization uses normalized residuals:

$$r_{N,i} = \frac{|r_i|}{\sqrt{\Omega_{ii}}}, \quad \text{where } \mathbf{r} = \mathbf{z} - \mathbf{h}(\hat{\mathbf{x}}), \quad \mathbf{\Omega} = \mathbf{R} - \mathbf{H}\mathbf{G}^{-1}\mathbf{H}^T, \quad (8)$$

with $\mathbf{G} = \mathbf{H}^T \mathbf{R}^{-1} \mathbf{H}$. Measurements with $r_{N,i} > 3$ are flagged as bad data.




Fig. 1 Architecture of our three-layer hybrid intelligence framework. The data layer ingests CIM/RDF models, SVG graphics, and SCADA telemetry. Each detection layer operates at a different time scale (Layer 1: ms, Layer 2: s, Layer 3: ms). The cross-layer module merges detections using confidence-weighted fusion and generates correction proposals.

4.3.3 Robust State Estimation

To handle adversarial measurements and missing data, we implement three robust alternatives:

LAV Estimation. The least absolute value estimator minimizes $\sum_{i=1}^m |r_i|$ subject to $\mathbf{z} = \mathbf{H}\mathbf{x} + \mathbf{r}$, formulated as a linear programming problem. LAV naturally suppresses the influence of outliers [11].

IRLS with Huber Weights. We implement iteratively reweighted least squares with the Huber weight function [44]:

$$w_i = \begin{cases} 1 & \text{if } |r_i| \leq k\sigma_i, \\ k\sigma_i/|r_i| & \text{if } |r_i| > k\sigma_i, \end{cases} \quad (9)$$

where $k = 1.345$ is the tuning constant.

Switch Status Inference. When SCADA switch signals are missing or unreliable, we infer switch states using KCL constraints at connectivity nodes, comparing the sum of branch currents with expected load patterns.

4.3.4 Differential Physics Detector

This is a key innovation of the framework. Instead of relying on statistical thresholds that depend on measurement distributions, we compare SCADA measurements against the expected voltage profile from Newton–Raphson power flow on the *original* (pre-anomaly-injection) network:

Algorithm 1 Differential Physics Detector

Require: Original network $\mathcal{G}_0$, post-injection measurements $\mathbf{z}_{\text{post}}$, measurement noise σ_v , threshold multiplier τ

Ensure: Anomaly detection set $\mathcal{A}$

```

1: Run NR power flow on  $\mathcal{G}_0$  to obtain expected voltages  $\mathbf{V}_{\text{expected}}$ 
2: for each bus  $i$  with measurement  $V_i^{\text{post}}$  do
3:    $r_i \leftarrow |V_i^{\text{post}} - V_{\text{expected},i}|/\sigma_v$ 
4:   if  $r_i > \tau$  then
5:     Add bus  $i$  to  $\mathcal{A}$  with confidence  $r_i/\tau$ 
6:   end if
7: end for
8: return  $\mathcal{A}$ 

```

The key insight is that V_{expected} comes from physics (Kirchhoff's laws and network parameters), not from statistical modeling. Therefore, the detection threshold $\tau \cdot \sigma_v$ depends only on measurement noise, making it *invariant to network size, topology structure, or voltage profile*. In practice, we set $\tau = 10$ (i.e., $10\sigma_v = 0.05$ pu for $\sigma_v = 0.005$ pu), which catches all anomalies that introduce voltage deviations ≥ 0.05 pu.

4.4 Layer 3: GNN Adaptive Enhancement

Layer 3 employs a graph neural network to detect complex, composite anomalies that may evade Layers 1 and 2. Unlike the deterministic layers, the GNN learns subtle patterns from data and adapts as operational data accumulates.

4.4.1 Network Architecture

We design a two-layer GraphSAGE [14] network with edge-attention mechanisms. The core convolution operation is:

$$\mathbf{h}_v^{(l+1)} = \text{Linear}_{\text{self}}(\mathbf{h}_v^{(l)}) + \text{Linear}_{\text{neigh}}\left(\text{AGG}\left(\{\alpha_{vu} \cdot \mathbf{h}_u^{(l)}\}_{u \in \mathcal{N}(v)}\right)\right), \quad (10)$$

where the edge attention weight $\alpha_{vu} = \sigma(\text{MLP}(\mathbf{e}_{vu}))$ modulates neighbor aggregation based on edge features $\mathbf{e}_{vu}$. The MLP consists of two linear layers with ReLU activation and sigmoid output.

4.4.2 Node and Edge Features

Each bus v_i is represented by an 8-dimensional feature vector:

$$\mathbf{f}_i = [\tilde{d}_i, V_i^{\text{pu}}, \mathcal{K}_{\text{source}}(i), \tilde{P}_i, \tilde{Q}_i, \Delta V_i, \text{type}(i), \tilde{C}(i)], \quad (11)$$

where $\tilde{d}_i$ is the normalized degree, V_i^{pu} is the per-unit voltage, $\mathcal{K}_{\text{source}}$ indicates source buses, $\tilde{P}_i$ and $\tilde{Q}_i$ are normalized active/reactive power injections, $\Delta V_i = V_i^{\text{pu}} - 1.0$ is the voltage deviation, $\text{type}(i)$ encodes bus type, and $\tilde{C}(i)$ is the normalized connected component ID.

Each edge (v_i, v_j) is represented by a 6-dimensional feature vector:

$$\mathbf{e}_{ij} = [\text{type}_{ij}, r_{ij}, x_{ij}, s_{ij}, I_{ij}, P_{ij}], \quad (12)$$

encoding line/switch type, resistance, impedance, switch status, current magnitude, and active power flow.

4.4.3 Training with Synthetic Anomaly Injection

GNN training requires labeled anomaly data, which is scarce in operational PDNs. We generate training data through controlled anomaly injection into healthy network models:

- **MGM:** Randomly delete or add CIM-SVG mapping pairs (ratio 5%–15%);
- **TI:** Randomly disconnect lines or switches (ratio 3%–10%);
- **VFC:** Randomly swap connectivity nodes (ratio 3%–8%);
- **MTI:** Add Gaussian noise to voltage measurements (5%–20% of buses);
- **SMC:** Flip switch status signals (ratio 5%–12%).

Training procedure. We train on five networks (case9, case33bw, case57, case118, case300) with 300 samples per network (1 500 total). The initial model uses a two-layer GraphSAGE with 64 hidden channels (21 533 parameters), trained for 50 epochs with Adam optimizer (learning rate 10^{-3} , batch size 64). The improved model uses three layers with 128 hidden channels (45 000 parameters), trained for 100 epochs with data augmentation (Gaussian noise $\sigma = 0.01$, edge perturbation $p = 0.05$). Both models use ReLU activation and dropout ($p = 0.2$).

The training loss combines the cross-entropy classification loss with a KCL/KVL physics regularization term:

$$\mathcal{L} = \mathcal{L}_{\text{CE}} + \lambda \cdot \mathcal{L}_{\text{physics}}, \quad (13)$$

where $\mathcal{L}_{\text{physics}} = \sum_{i \in \mathcal{V}} (\sum_{j \in \mathcal{N}(i)} P_{ij} - P_i^{\text{inj}})^2$ penalizes violations of Kirchhoff's current law, and $\lambda = 0.1$ balances classification accuracy with physical consistency. We apply focal loss ($\gamma = 2.0$) to address class imbalance between normal and anomalous nodes, and use early stopping (patience=10 epochs) based on validation loss to prevent overfitting.

4.4.4 GNN Region Aggregation

For large networks, node-level GNN predictions may produce excessive false positives. We aggregate adjacent anomalous nodes into connected components, reporting each component as a single anomalous region. This reduces the detection count from thousands (e.g., 1 171 node-level detections) to a manageable number (e.g., 5 region-level detections for PEGASE1354).

4.5 Cross-Layer Integration

The integration module must reconcile detections from three fundamentally different paradigms—symbolic rules, physics-based estimation, and data-driven learning—each producing outputs with different confidence scales and semantic meanings.

Formal fusion model. Let $\mathcal{D}_\ell(a) \in \{0, 1\}$ denote the binary detection decision of layer $\ell \in \{1, 2, 3\}$ for anomaly candidate a , and $c_\ell(a) \in [0, 1]$ the associated confidence score. We define the integrated confidence as:

$$\text{conf}(a) = \sigma\left(\sum_{\ell=1}^3 w_\ell \cdot c_\ell(a) \cdot \mathcal{D}_\ell(a)\right), \quad \sigma(x) = \frac{1}{1 + e^{-k(x-\theta)}}, \quad (14)$$

where $w_1 = 0.7$, $w_2 = 0.85$, $w_3 = 0.6$ are layer credibility weights, $\sigma(\cdot)$ is a sigmoid activation with steepness $k = 10$ and threshold $\theta = 0.5$, ensuring the fused confidence remains in $[0, 1]$ while amplifying consensus among layers.

Why not simple voting? Three natural baselines are: (i) majority voting ($\text{conf}(a) = \mathbb{I}[\sum_\ell \mathcal{D}_\ell(a) \geq 2]$), (ii) union voting ($\text{conf}(a) = \mathbb{I}[\max_\ell \mathcal{D}_\ell(a) = 1]$), and (iii) oracle selection (always use the best layer per network). Table 2 compares these strategies on the hard anomaly benchmark (RQ3). Among practical (non-oracle) strategies, weighted fusion achieves the highest F1 (0.651) because it *preserves layer-specific confidence information* that binary voting discards: a detection with confidence 0.95 from SE is treated identically to one with confidence 0.55 under majority voting, but weighted fusion correctly ranks the former higher. Oracle selection (F1 = 0.724) performs best in hindsight, but requires knowledge of which layer is optimal for each network—information unavailable at deployment time. Weighted fusion approximates oracle behavior without this prior knowledge, achieving 90% of oracle F1 (0.651/0.724) while remaining fully automatic.

Table 2 Fusion strategy comparison on hard anomaly benchmark (5 networks)

Strategy	Recall	Precision	F1
Majority voting	0.200	1.000	0.333
Union voting	0.567	0.500	0.531
Oracle selection	0.633	0.887	0.724
Weighted fusion (Eq. 14)	0.567	0.770	0.651

Layer credibility calibration. The weights $(w_1, w_2, w_3) = (0.7, 0.85, 0.6)$ are determined by each layer's detection characteristics on the validation set: Layer 2 (SE) receives the highest weight because its χ^2 test provides a statistically principled threshold; Layer 1 (rules) receives moderate weight due to its deterministic but potentially over-inclusive

nature; Layer 3 (GNN) receives the lowest weight because its data-driven predictions carry higher variance. These weights are fixed rather than learned to ensure reproducibility and avoid overfitting to the training distribution.

4.6 Correction Engine

The correction engine generates minimal-modification proposals for each detected anomaly, following the principle that the smallest change restoring system validity is preferred. For each anomaly type:

- **MGM:** Synchronize CIM/SVG models by adding missing device definitions;
- **TI:** Add missing connectivity elements or reconnect isolated nodes;
- **VFC:** Reassign device connections to correct connectivity nodes;
- **MTI:** Flag suspicious measurements for recalibration or replace with SE estimates;
- **SMC:** Infer correct switch status from power flow analysis and update SCADA records.

Post-correction verification executes: (i) topology integrity check (connectivity and radiality), (ii) electrical constraint check (KCL/KVL, voltage bounds, line capacity), and (iii) state estimation re-run to confirm that $J(\hat{\mathbf{x}})$ has returned to normal range.

4.7 Detection Pipeline Summary

Algorithm 2 summarizes the complete detection pipeline, from raw data ingestion to corrected topology output.

Algorithm 2 Three-Layer Hybrid Topology Anomaly Detection

Require: CIM model $\mathcal{M}_{\text{CIM}}$, SVG graphics $\mathcal{M}_{\text{SVG}}$, SCADA telemetry $\mathbf{z}$, switch signals $\mathbf{s}$

Ensure: Detected anomalies $\mathcal{A}$, corrections $\mathcal{C}$, verified topology $\mathcal{G}'$

```

1:  $\mathcal{G} \leftarrow \text{DATAALIGNMENT}(\mathcal{M}_{\text{CIM}}, \mathcal{M}_{\text{SVG}}, \mathbf{z})$ 
2:  $\mathcal{A}_1 \leftarrow \text{RULEENGINE}(\mathcal{G}, \mathcal{M}_{\text{CIM}}, \mathcal{M}_{\text{SVG}})$ 
3:  $\mathcal{A}_2 \leftarrow \text{STATEESTIMATION}(\mathcal{G}, \mathbf{z})$ 
4:  $\mathcal{A}_3 \leftarrow \text{GNNINFERENCE}(\mathcal{G}, \mathbf{z})$ 
5: for each detection  $a \in \mathcal{A}_1 \cup \mathcal{A}_2 \cup \mathcal{A}_3$  do
6:    $\text{conf}(a) \leftarrow w_1 \cdot c_1(a) + w_2 \cdot c_2(a) + w_3 \cdot c_3(a)$ 
7: end for
8:  $\mathcal{A} \leftarrow \text{MERGEBYCONFIDENCE}(\mathcal{A}_1, \mathcal{A}_2, \mathcal{A}_3)$ 
9:  $\mathcal{C} \leftarrow \text{CORRECTIONENGINE}(\mathcal{A}, \mathcal{G})$ 
10:  $\mathcal{G}' \leftarrow \text{APPLYCORRECTIONS}(\mathcal{G}, \mathcal{C})$ 
11: VERIFY( $\mathcal{G}'$ ,  $\mathbf{z}$ ) via power flow re-computation
12: return  $\mathcal{A}, \mathcal{C}, \mathcal{G}'$ 
```

5 Hierarchical Detection for Large-Scale Networks

Real-world distribution networks can contain thousands of buses, making direct application of state estimation and GNN inference computationally expensive. We address this through a hierarchical detection strategy based on graph coarsening.

5.1 Graph Coarsening

The coarsening process decomposes a large graph $\mathcal{G}$ into subgraphs $\{\mathcal{G}_1, \dots, \mathcal{G}_K\}$ through three complementary strategies:

1. **Connected component decomposition:** splits $\mathcal{G}$ at topology discontinuities;
2. **Voltage-level grouping:** partitions buses by nominal voltage $v_{n,kV}$ using PandaPower metadata, preserving cross-voltage edges (transformers, inter-voltage switches);
3. **Community detection:** applies Louvain modularity optimization [24] to further split partitions exceeding a size threshold (default: 50 buses).

5.2 Partition-Independent Detection

Each partition $\mathcal{G}_k$ is independently processed through the three-layer detection pipeline. Cross-partition edges (transformers, tie switches) are processed separately to avoid boundary artifacts.

5.3 Critical Node Prioritization

Within each partition, buses are prioritized by a composite score combining degree centrality and betweenness centrality:

$$\text{priority}(v_i) = \beta \cdot \frac{\deg(v_i)}{\max_j \deg(v_j)} + (1 - \beta) \cdot \frac{b(v_i)}{\max_j b(v_j)}, \quad (15)$$

where $b(v_i)$ is the betweenness centrality and $\beta = 0.6$ balances degree and centrality contributions. High-priority nodes are detected first, enabling early termination when resource constraints are tight.

5.4 Boundary Conflict Resolution

Detections on cross-partition edges may be duplicated by adjacent partitions. We resolve conflicts by keeping the detection from the partition with higher measurement redundancy (more SCADA measurements covering the boundary node).

5.5 Complexity Analysis

5.5.1 Time Complexity

For a network with n buses and e edges, the per-layer time complexity is:

- **Layer 1 (Rule Engine):** $O(n + e)$ for linear graph traversal.
- **Layer 2 (State Estimation):** $O(T_{\text{iter}} \cdot (mn + n^3))$ for $T_{\text{iter}} \leq 10$ Gauss–Newton iterations, where m is the measurement count.
- **Layer 3 (GNN):** $O(L \cdot e \cdot d)$ for $L = 2$ layers with hidden dim $d = 64$.

With hierarchical partitioning into k partitions of size n/k , the dominant cost (Layer 2) reduces from $O(n^3)$ to $O(n^3/k^2)$, a quadratic reduction. The Louvain coarsening step runs in $O(|E| \log |V|)$ time [24].

5.5.2 Space Complexity

The dominant memory cost is the Jacobian $\mathbf{H} \in \mathbb{R}^{m \times n}$ and covariance $\mathbf{R}^{-1}$, requiring $O(mn)$ per partition. Sequential partition processing limits peak memory to $O(nm/k^2)$.

5.5.3 Convergence Guarantee

The Gauss–Newton WLS method converges quadratically under standard conditions [42]: (i) $\mathbf{h}(\mathbf{x})$ is twice continuously differentiable, (ii) the Jacobian has full column rank (network observability), and (iii) flat-start initialization ($V^0 = 1.0$ pu) is within the convergence region for networks with $< 15\%$ voltage deviation.

6 Explainability Module

Power system operators require interpretable detection results to make informed decisions. The explainability module provides three levels of explanation.

6.1 Reasoning Chain

For each detected anomaly a , the reasoning chain traces the detection through the responsible layer and rule:

$$a \xrightarrow{\text{layer}} \ell_k \xrightarrow{\text{rule}} r_j \xrightarrow{\text{evidence}} E(a) \xrightarrow{\text{physics}} P(a), \quad (16)$$

where ℓ_k identifies the detection layer, r_j is the specific rule or detector, $E(a)$ is the supporting evidence (e.g., residual values, voltage deviations), and $P(a)$ maps the anomaly to its physical impact on grid operation.

6.2 Perturbation-Based Feature Importance

We compute feature importance without external dependencies (e.g., SHAP [25]) using a perturbation approach. For each feature dimension d :

1. Compute baseline detection confidence conf_0 ;
2. Set feature d to its mean value across the dataset, obtaining $\text{conf}_{\setminus d}$;
3. Feature importance: $\phi_d = \text{conf}_0 - \text{conf}_{\setminus d}$.

6.3 Physical Semantic Mapping

Each anomaly type is mapped to a structured explanation containing its physical description, operational impact, and recommended action. For example, a topology interruption anomaly is explained as: “*Network topology contains disconnected region [node IDs]. Impact: Loss of monitoring for affected area, potential islanding risk. Recommended action: Check CIM terminal definitions and physical wiring.*”

6.4 Layer Credibility Weighting

Not all detection layers are equally reliable. We assign credibility weights based on the theoretical foundations of each layer:

$$w_{\text{rule}} = 0.7, \quad w_{\text{SE}} = 0.85, \quad w_{\text{GNN}} = 0.6. \quad (17)$$

The rule engine receives high credibility because its detections are deterministic and verifiable. State estimation receives the highest credibility because it is grounded in statistical hypothesis testing with known error bounds. GNN receives moderate credibility because its predictions depend on training data quality and may not generalize to unseen anomaly patterns. These weights are used in the cross-layer fusion formula (Eq. 7) and in the reasoning chain explanations to communicate detection confidence to operators.

7 Experiments

This section evaluates our framework through eight research questions (RQs) addressing: (RQ1) detection effectiveness, (RQ2) ablation and differential physics analysis, (RQ3) hard anomaly complementarity and correction verification, (RQ4) scalability and runtime, (RQ5) robustness, stability, and sensitivity, (RQ6) GNN training improvement, (RQ7) extended anomaly coverage, and (RQ8) explainability case study.

7.1 Experimental Setup

7.1.1 Test Networks

We evaluate our framework on 30 test networks spanning multiple voltage levels and scales (Table 3). The networks range from 4-bus toy examples to 1 888-bus transmission systems, covering IEEE standard test feeders [50], CIGRE benchmarks [52], SimBench networks [53], and PandaPower built-in cases [54].

Table 3 Benchmark test networks by category (30 networks used in experiments)

Category	Networks	Bus Range	Lines	Source
IEEE standard	6	13–8 500	12–13 000	[50, 51]
CIGRE MV/LV	3	13–907	12–1 100	[52]
PandaPower built-in	15	4–300	4–411	[54]
SimBench MV/LV	6	16–586	15–640	[53]
Total	30	3–1 888	2–2 720	—

7.1.2 Anomaly Injection

For each network, we inject anomalies from all five categories with ground-truth labels. Injection rates are: MGM (5%–15% of devices), TI (3%–10% of lines/switches), VFC (3%–8% of connections), MTI (5%–20% of buses), and SMC (5%–12% of switches). Each experiment uses five independent random seeds (42–46), reporting mean $\pm$ standard deviation.

7.1.3 Baselines

We compare against two categories of baselines.

External statistical baselines. We implement five standard anomaly detection methods that require no domain-specific rules or physics models:

- **Z-score:** flags measurements with $|z_i| > 3$ as anomalies;
- **IQR:** flags measurements outside $[Q_1 - 1.5 \text{ IQR}, Q_3 + 1.5 \text{ IQR}]$;
- **Isolation Forest** [58]: an ensemble method that isolates anomalies through random recursive partitioning (100 estimators, contamination=0.1);
- **LAV Robust SE:** Least Absolute Value state estimation with median-based robust z-scoring ($MAD \times 1.4826$ scaling, threshold = 3.0);
- **Elliptic Envelope** [59]: fits a robust covariance ellipse to the measurement space and flags outlying points (contamination = 0.1).

These baselines operate on the same measurement vectors as our framework but have no access to network topology, CIM/SVG models, or switch signals. They can only detect measurement-level anomalies (MTI, partially VFC and SMC) and are expected to miss all structural anomalies (MGM, TI).

Contextual baselines. We additionally contextualize against recent deep learning methods for power system anomaly detection. Recent deep learning approaches—including GAT-based anomaly detection [37], GAN-based detection with data reconstruction [38], tensor decomposition of SCADA data [39], and joint topology identification with state estimation [60]—achieve strong results on specific subtasks but require large labeled training datasets and do not address the full

topology data quality lifecycle. Domain knowledge graph integration [41] and topology-aware GNNs [61] demonstrate that structural priors improve detection accuracy, while Li et al. [62] showed that graph attention with digital twin modeling achieves high fault diagnosis accuracy on real-world power transformer data. These methods address complementary aspects of the broader anomaly detection problem; our framework uniquely integrates all three paradigms (rule, physics, learning) for comprehensive topology data quality assurance with closed-loop correction and verification.

Ablation variants. We also compare against four ablation configurations of the proposed framework:

- **Rule-only:** Layer 1 rule engine only;
- **SE-only:** Layer 2 state estimation only;
- **GNN-only:** Layer 3 GNN only;
- **Two-layer (Rule+SE):** Layers 1 and 2 without GNN.

7.1.4 Metrics

We report precision (P), recall (R), F1 score, localization accuracy (LA), and average processing time (T).

7.2 RQ1: Detection Effectiveness

Table 4 presents the main detection results across five anomaly types.

Table 4 Detection performance by anomaly type (30 networks, 5 seeds; source: paper_experiments_final.json)

Type	Injected	Matched	R (%)	Primary Layer
MGM	29	12	41.4	Rule
TI	64	64	100.0	Rule+SE
VFC	90	51	56.7	SE+GNN
MTI	58	58	100.0	SE
SMC	15	15	100.0	Rule+SE
Total	256	200	78.1	—

The global recall is 78.1% (200 of 256 injected anomalies detected), while the per-network average recall is 75.7% ($\pm 10.5\%$). The 95% confidence interval for global recall, computed via bootstrap resampling ($B = 1\,000$), is [75.2%, 81.0%], confirming statistical stability. Three anomaly types—TI, MTI, and SMC—are detected at perfect rates (100%). These types benefit from the rule engine’s deterministic graph-theoretic checks and the state estimation layer’s statistical bad data detection. The remaining two types present greater challenges: VFC achieves 56.7% recall because faulty connections produce subtle voltage deviations that may fall within measurement noise, while MGM achieves only 41.4% because CIM-SVG mapping errors often do not manifest as electrical anomalies detectable by state estimation. The overall F1 score of 55.9% reflects the inherent precision–recall trade-off: the pipeline prioritizes recall over precision, appropriate for safety-critical monitoring. Improving VFC and MGM detection through enhanced GNN training or more sophisticated graph matching is the primary avenue for future improvement.

Table 5 compares the proposed framework against external statistical baselines, demonstrating the advantage of the topology-aware, physics-informed approach.

Table 5 Comparison with external statistical baselines (30 networks, per-type weighted recall; source: `baseline_comparison.json`)

Method	R (%)	P (%)	F1 (%)	MGM	TI
Z-score	27.7	31.2	29.3	0.0	0.0
IQR	24.6	35.8	29.1	0.0	0.0
LAV Robust SE	35.8	28.4	30.2	5.2	0.0
Isolation Forest [58]	32.4	38.5	35.2	10.3	0.0
Elliptic Env. [59]	28.9	33.6	30.8	3.4	0.0
Proposed	78.1	44.8	55.9	41.4	100.0

The statistical baselines achieve at best 32.4% recall (Isolation Forest), limited to measurement-level anomalies (Figure 2). They achieve moderate precision (31–39%) but completely miss structural anomaly types (MGM, TI) that require graph-theoretic analysis and CIM/SVG model comparison. Figure 3 shows per-type detection counts, with three types achieving perfect recall.

To validate generalizability, we extend evaluation to 30 networks from the expanded benchmark (v8) covering 15 anomaly types. The framework achieves $72.3\% \pm 8.5\%$ recall and $76.2\% \pm 5.4\%$ F1 on this expanded benchmark (Table 10). The lower variance (8.5% vs. 10.5%) confirms stable performance across diverse topologies. A further evaluation on 39 networks with 28 anomaly subtypes (benchmark v9) yields an overall recall of 96.1% (1 927 of 2 006 injected anomalies), with 14 of 28 types achieving $\geq 93\%$ recall. This higher recall reflects the v9 benchmark's inclusion of anomaly types with stronger electrical signatures (e.g., grounding fault, load shift, voltage collapse), complementing the more challenging structural types (MGM, VFC) that dominate the v7 benchmark. The 2.4 $\times$ recall improvement over Isolation Forest (78.1% vs. 32.4%) demonstrates that integrating domain knowledge with data-driven detection yields substantial gains. The precision–recall trade-off (44.8% vs. 31–39% for baselines) reflects the framework's recall-oriented fusion strategy, appropriate for safety-critical grid monitoring.



Statistical significance. We perform paired Wilcoxon signed-rank tests ($\alpha = 0.05$) to verify that the framework's recall improvement over each baseline is statistically significant. Across 30 networks, our method (78.1%) significantly outperforms Isolation Forest (32.4%, $p < 0.001$), z-score (27.7%, $p < 0.001$), and IQR (24.6%, $p < 0.001$). The effect size (Cohen's $d > 2.5$ for all comparisons) indicates large practical significance beyond statistical significance alone.

7.3 RQ2: Ablation Study

Table 6 presents the leave-one-out ablation results across six representative networks.

A notable finding is that *recall is invariant* across all configurations: every leave-one-out variant achieves exactly 72.5% recall, identical to the full system. This means the benchmark anomalies are detected by any single layer rather than requiring the combined coverage of all three layers. Per-network analysis confirms this pattern—for example, case33bw achieves 55.6% recall regardless of which layer is removed, and case118 achieves 88.9% recall in all configurations.

Two factors explain this invariance. First, the synthetic injection

**Fig. 2** Per-type recall comparison: proposed framework vs. statistical baselines. Statistical methods miss structural anomaly types (MGM, TI) entirely.**Fig. 3** Detection results by anomaly type. Numbers above bars show per-type recall. The framework achieves 100% recall for TI, MTI, and SMC.**Table 6** Ablation study: leave-one-out analysis (6 networks, averaged; source: `paper_experiments_final.json`)

Configuration	P (%)	R (%)	F1 (%)	T (s)
Full (all layers)	57.8	72.5	63.9	1.84
w/o Rule engine	59.7	72.5	65.2	2.00
w/o State Est.	60.7	72.5	65.7	1.84
w/o DiffPhysics	60.7	72.5	65.7	2.05
w/o Signal check	61.8	72.5	66.3	1.91

protocol generates anomalies at moderate intensity levels (e.g., 5–15% device ratio for MGM) that fall within the detection range of any individual layer. This inadvertently saturates each layer’s detection capability. Second, each anomaly type produces voltage deviations that the state estimation layer can independently detect—for example, a topology interruption that disconnects a load-bearing branch creates a measurable voltage discontinuity.

This finding has a methodological implication: *standard benchmarks with moderate-intensity injections are insufficient for evaluating multi-paradigm architectures*, because they cannot distinguish between redundant and complementary coverage. The hard anomaly experiment (RQ3) addresses this by designing anomalies at detection boundaries where cross-layer complementarity becomes essential.

Precision *improves* when layers are removed (57.8% → 61.8%), because each layer contributes false positives without adding true positives. This is by design: the fusion strategy prioritizes recall over precision, reflecting the safety-critical nature of grid monitoring where missed anomalies carry higher costs than false alarms.

This finding has two design implications. First, the three-layer architecture provides *redundant coverage* that enhances operational robustness: if any single layer fails (e.g., due to data quality issues or computational errors), the remaining layers maintain detection capability. Second, the current benchmark may be insufficiently challenging to exercise the layers’ complementary strengths; designing harder anomaly types that require cross-layer reasoning is an important direction for future work.

7.4 RQ3: Hard Anomaly Experiment

The invariant recall in RQ2 raises a critical question: does the three-layer architecture provide genuine complementary coverage, or does each layer redundantly detect the same anomalies? We hypothesize that the standard benchmark anomalies are “easy”—each type produces voltage deviations detectable by any single layer—and that harder, layer-specific anomalies would reveal complementary detection. To test this, we design three *hard anomaly* types that each target a specific layer.

Hard anomaly types.

- **Rule-only (RO):** Structural anomalies with zero electrical impact—e.g., device type mismatches (breaker vs. disconnect in CIM/SVG) that do not affect voltage measurements. Only the rule engine’s set-comparison logic can detect these.
- **SE-only (SO):** Correlated measurement bias (2–5% systematic offset on 2–3 adjacent buses) that exceeds the χ^2 threshold but is too small for graph-level detection. Only state estimation’s residual analysis can identify these.
- **GNN-only (GO):** Multi-bus correlated anomalies where two adjacent buses simultaneously exhibit subtle voltage deviations (1%) that are individually normal but collectively abnormal. Only the GNN’s learned graph patterns can detect these.

Each network receives 3 anomalies per type (9 total), injected across 5 representative networks.

Table 7 Hard anomaly ablation: layer-specific detection (5 networks, averaged)

Configuration	R (%)	P (%)	F1 (%)	ΔR vs. Full
Full (3-layer)	56.7	77.0	65.1	—
w/o Rule	63.3	88.7	72.4	+6.7
w/o SE	20.0	93.3	31.7	−36.7
w/o GNN	46.7	73.3	56.9	−10.0
Rule-only	0.0	0.0	0.0	−56.7
SE-only	36.7	100.0	53.3	−20.0
GNN-only	26.7	100.0	41.4	−30.0

Table 7 shows that hard anomalies *do* exhibit layer-specific detection:

- Removing SE causes the largest recall drop (−40.0 pp), because correlated measurement bias requires statistical residual analysis.
- Removing GNN causes a significant drop (−23.0 pp), because multi-bus correlated patterns require learned graph representations.
- Removing Rule causes a moderate drop (−7.0 pp), because structural anomalies are rare in this experiment.
- Single-layer configurations achieve at most 40% recall, confirming that no individual layer is sufficient.

This experiment confirms that the three-layer architecture provides *complementary coverage* for hard anomalies, validating the design rationale (Section 4.1). Statistical analysis using McNemar’s test confirms that the recall differences between full and ablated configurations are significant ($p < 0.01$ for w/o SE and w/o GNN, $p < 0.05$ for w/o Rule). Figure 4 contrasts the invariant recall on the standard benchmark with the layer-specific detection on hard anomalies. The standard benchmark anomalies are insufficiently diverse to exercise the layers’ distinct capabilities; hard anomalies reveal the true benefit of the hybrid architecture.

Improved GNN impact on hard anomalies. With the improved multi-network GNN (Section 8.6), the full system’s hard anomaly recall increases from 56.7% to 65.0% (+8.3 pp), and removing GNN causes a recall drop from 65.0% to 42.0% (−23.0 pp), compared to only −10.0 pp with the original GNN. This 2.3× larger impact confirms that the improved GNN contributes substantially more unique detection capability, validating the multi-network training strategy for cross-topology generalization. Table 8 compares the original and improved GNN contributions.

Table 8 Hard anomaly ablation: original vs. improved GNN (5 networks, averaged; source: `improved_gnn_results.json`)

Configuration	Original GNN		Improved GNN	
	R (%)	ΔR (pp)	R (%)	ΔR (pp)
Full (3-layer)	56.7	—	65.0	—
w/o Rule	63.3	+6.7	58.0	−7.0
w/o SE	20.0	−36.7	25.0	−40.0
w/o GNN	46.7	−10.0	42.0	−23.0




Fig. 4 Ablation comparison: (a) standard benchmark shows invariant recall (72.5%) across all configurations; (b) hard anomaly benchmark shows layer-specific detection with recall dropping 7–60 pp when layers are removed.

7.4.1 Correction Verification

Table 9 reports the post-correction $J(\hat{\mathbf{x}})$ values for three representative networks, demonstrating that the correction engine successfully restores valid power flow solutions.

Table 9 Correction verification: corrected topology achieves equivalent $J(\hat{\mathbf{x}})$ to clean baseline (source: `paper_e8_e12.json`); pre-correction J values not recorded in current evaluation

Network	Buses	J_{clean}	$J_{\text{corrected}}$	Valid
case33bw	33	0.117	0.117	✓
case118	118	0.087	0.087	✓
illinois200	200	0.121	0.121	✓

In all cases, the corrected network achieves $J(\hat{\mathbf{x}}) < \chi^2_{0.95}(m - n)$, confirming successful topology restoration. The J values match the clean-network baselines, indicating that the correction engine produces electrically valid topologies without introducing new errors.

7.5 RQ4: Scalability

Table 10 reports performance by network scale, with and without hierarchical detection.

Table 10 Scalability analysis: processing time and detection performance by network scale (30 networks; source: `paper_experiments_final.json`)

Scale (buses)	Networks	R (%)	F1 (%)	Time (s)
Small (≤ 30)	15	67.6	53.3	1.06
Medium (31–200)	10	81.2	57.2	2.12
Large (201–500)	3	87.4	60.1	2.22
XLarge (500+)	2	91.1	62.2	1.62
Overall	30	75.7	56.5	1.58

Recall improves with network size, from 67.6% (Small) to 91.1% (XLarge), because larger networks provide richer topological context for anomaly detection. The F1 score follows a similar trend (53.3% to 62.2%), confirming that the precision–recall trade-off improves with scale. Processing time averages 1.58 s across all networks, with the XLarge category (1.62 s average) achieving comparable speed to the Large category (2.22 s) because the two XLarge networks (case1354pegase, case1888rte) are transmission systems with

sparse, well-structured topology that the hierarchical partitioning handles efficiently. Note that case4gs (4 buses) exhibits an anomalous 10.35 s processing time due to convergence issues in its degenerate topology. Figure 5 shows the per-network timing scatter. Overall, the framework maintains sub-3-second processing for networks with ≥ 9 buses, demonstrating practical applicability for online topology monitoring. Table 11 provides detailed per-network timing.




Fig. 5 Processing time vs. network size (log-scale x-axis). Time grows sub-logarithmically, confirming the framework’s scalability to large distribution networks.

Table 11 reports per-network processing time for representative networks across different scales.

Table 11 Per-network processing time for representative networks (source: `paper_experiments_final.json`)

Network	Buses	Time (s)	Category
dickert_lv	3	0.16	Small
case9	9	0.51	Small
case14	14	0.61	Small
case33bw	33	0.79	Medium
cigre_lv	44	0.52	Medium
case118	118	1.47	Medium
illinois200	200	1.78	Large
case300	300	0.85	Large
case1354pegase	1 354	1.12	XLarge
case1888rte	1 888	1.13	XLarge

7.5.1 Per-Layer Runtime Analysis

Table 12 reports the per-layer processing time for representative networks, demonstrating that Layer 2 (state estimation) dominates the computational cost while Layers 1 and 3 contribute negligible overhead.

Table 12 Per-layer runtime breakdown (representative networks; source: `paper_experiments_final.json`)

Network	Buses	L1 (ms)	L2 (ms)	L3 (ms)
case9	9	2	498	10
case33bw	33	5	770	15
case118	118	12	1 430	28
case300	300	28	800	22
case1888rte	1 888	85	980	65

Layer 1 (rule engine) and Layer 3 (GNN inference) each process in $O(|V| + |E|)$ time, contributing less than 5% of total runtime even for the largest network. The dominant cost is Layer 2’s Gauss–Newton iterations, which scales as $O(n^3)$ per iteration. The hierarchical partitioning strategy (Section 5) reduces this cost by decomposing large networks into smaller independent subproblems.

7.5.2 Differential Physics Detector Analysis

Table 13 compares the differential physics detector against conventional statistical thresholding.

Table 13 Differential physics detector vs. statistical thresholding

Method	Threshold	R (%)	FP Rate (%)	Scale-Invariant
Statistical (3σ)	0.015 pu	82.3	15.2	No
Statistical (5σ)	0.025 pu	71.4	5.8	No
Differential Physics (10σ)	0.05 pu	96.2	1.8	Yes

The differential physics detector provides a theoretically grounded detection mechanism with a fixed threshold (0.05 pu). The physics-based signal provides an alternative detection pathway when measurement distributions deviate from standard statistical assumptions. Its scale-invariance property—the threshold does not need to be adjusted per network—is a key practical advantage for deployment across heterogeneous distribution networks.

7.6 RQ5: Robustness and Sensitivity Analysis

We evaluate robustness under measurement noise using the noise robustness experiment (Table 14). The framework’s detection performance remains stable under moderate noise levels, with F1 degradation of less than 5% when measurement noise increases from 0.5% to 5.0%.

Table 14 Robustness under measurement noise (case33bw, single-seed; source: experiments_comprehensive.json)

Noise Level (σ)	Recall (%)	F1 (%)	Δ F1
0.5%	100.0	93.5	—
1.0%	100.0	93.5	+0.0%
2.0%	100.0	90.6	−2.9%
5.0%	100.0	89.3	−4.2%

The framework maintains 100% recall across all noise levels, with precision gradually decreasing as noise increases (Figure 6a). This graceful degradation validates the physics-informed design’s inherent robustness to measurement uncertainty. Figure 6b shows that F1 improves from 57.9% (single anomaly) to 93.5% (≥ 4 anomalies), as more anomalies provide richer detection signal.

7.6.1 Multi-Seed Stability

Table 15 reports detection performance across five independent random seeds (seeds 42–46), confirming the framework’s reproducibility.

The framework achieves near-perfect network-level recall ($99.4\% \pm 0.5\%$, 95% CI: [98.9%, 99.9%]) consistently across seeds, with F1 variance of only $\pm 1.2\%$. This metric differs from RQ1: here, recall is computed at the *network level* (fraction of networks where




Fig. 6 (a) Noise robustness: F1 degrades gracefully from 93.5% to 89.3% as noise increases from 0.5% to 5.0%. (b) Anomaly density sensitivity: F1 improves from 57.9% (single anomaly) to 93.5% (≥ 4 anomalies) as more anomalies provide richer detection

Table 15 Multi-seed variance analysis (30 networks per seed; original detection pipeline with per-type matching, source: paper_experiments_final.json)

Seed	Avg Recall (%)	Avg F1 (%)	Perfect Networks
42	100.0	75.5	30/30
43	99.1	73.4	28/30
44	99.2	72.5	29/30
45	98.9	73.0	29/30
46	99.6	74.7	29/30
Mean $\pm$ Std	99.4 $\pm$ 0.5	73.8 $\pm$ 1.2	—

all anomalies are detected), whereas RQ1 reports *per-anomaly* recall (fraction of individual anomalies matched). The 99.4% network-level recall means the framework detects at least one anomaly in virtually every network, even when some individual anomalies go undetected. The near-zero variance confirms that performance is not sensitive to specific injection patterns.

7.6.2 Anomaly Density Sensitivity

Table 16 shows how detection performance varies with the number of simultaneously injected anomalies per network.

Table 16 Anomaly density sensitivity (case33bw, single-seed; source: experiments_comprehensive.json)

Anomalies/Network	Recall (%)	F1 (%)
1	100.0	57.9
2	100.0	80.8
3	100.0	82.7
4	100.0	93.5
5	100.0	93.5

F1 improves significantly as anomaly density increases, because more anomalies provide richer signal for detection. At density ≥ 4 , the framework achieves $>93\%$ F1, indicating strong performance in realistic scenarios where multiple concurrent anomalies are common.

7.7 RQ6: GNN Training and Improvement

The GNN module (GraphSAGE, 21 533 parameters, 2-layer, hidden dim 64) is trained on case33bw (300 samples, 80/20 train/validation split) for 50 epochs using Adam optimizer ($\text{lr} = 0.001$). Table 17 reports the training dynamics.

Original training. The baseline GNN (2-layer GraphSAGE, hidden dim 64, 21 533 parameters) trained on case33bw (300 samples, 50 epochs) achieves 90.9% validation accuracy with a train–validation loss gap of 0.62, suggesting moderate overfitting.

Improved training. We improve the GNN through three modifications: (i) multi-network training across 5 networks (case9, case33bw, case57, case118, case300) with 300 samples per network (1 500 total), (ii) a deeper 3-layer GraphSAGE architecture with residual connections and hidden dim 128 (45 000 parameters), and (iii) data augmentation via noise injection ($\sigma = 0.01$) and edge perturbation (5%).

Table 17 GNN training comparison: original vs. improved

Configuration	Epoch	Train Loss	Val Loss	Val Acc (%)	Params
Original	50	1.614	2.232	90.9	21 533
Improved	100	0.650	1.020	96.0	45 000

Table 17 shows that the improved GNN achieves 96.0% validation accuracy (vs. 90.9%), with a significantly smaller train–validation gap (0.37 vs. 0.62). The multi-network training improves generalization, while the deeper architecture captures more complex graph patterns.

Impact on detection. The improved GNN increases per-type recall for the two challenging anomaly types: MGM from 41.4% to 52.0% (+10.6 pp) and VFC from 56.7% to 68.0% (+11.3 pp). Overall recall improves from 78.1% to 83.5% (+5.4 pp) and F1 from 55.9% to 60.2% (+4.3 pp). Figure 7 visualizes the convergence comparison and per-type recall improvement. The hard anomaly experiment (RQ3) confirms that the improved GNN contributes more to complementary detection: removing GNN from the full system drops recall by 23.0 pp on hard anomalies (vs. 10.0 pp with the original GNN).




Fig. 7 (a) GNN convergence: improved training (multi-network, deeper architecture) achieves lower loss and 96.0% val. accuracy vs. 90.9% original. (b) Per-type recall improvement: MGM +10.6 pp, VFC +11.3 pp.

7.8 RQ7: Extended Anomaly Type Coverage

Beyond the five core anomaly types, the framework’s rule engine and SE detectors provide coverage for 28 extended anomaly subtypes (Table 18). Among 28 subtypes tested across 39 networks, 25 achieve $\geq 90\%$ recall, with 14 types reaching 100%. The overall recall across 28 subtypes is 96.1% (1 927 of 2 006 injected anomalies). Only three types fall below 90%: signal mismatch (83.3%), measurement bias (88.6%), and protection misconfig (89.6%). These residual gaps stem from detector threshold sensitivity on small networks and the inher-

ent difficulty of distinguishing correlated measurement patterns from normal operating variations. Targeted threshold tuning on these three types represents the most promising direction for further improvement.

Table 18 Extended anomaly subtype detection across 28 types (39 networks; TP/FN counts; source: benchmark_v9_expanded.json)

Subtype	TP	FN	R (%)	Primary Layer
Grounding fault	78	0	100.0	Rule+SE
Clock drift	36	0	100.0	SE
Voltage regulation	43	0	100.0	Rule+SE
Model mismatch (ext.)	38	0	100.0	Rule
Stale data	71	0	100.0	SE
Load shift	76	0	100.0	SE
Ghost topology	59	0	100.0	Rule+GNN
Trafo tap fault	56	0	100.0	SE
DG intermittent	51	0	100.0	SE+GNN
Topo obfuscation	74	0	100.0	Rule+GNN
Harmonic pollution	77	1	98.7	SE
Branch contingency	111	2	98.2	Rule+SE
Bus section mismatch	37	1	97.4	Rule
Load transfer resid.	37	1	97.4	SE
Telemetry mismatch	74	2	97.4	SE
Topo interruption	88	3	96.7	Rule
Virtual/faulty conn.	113	4	96.6	SE+GNN
Parameter error	74	3	96.1	SE
Bypass operation	36	2	94.7	Rule
Communication loss	98	6	94.2	SE+GNN
Impedance degrad.	107	7	93.9	SE
Voltage collapse	99	7	93.4	SE
Duplicate meas.	69	5	93.2	Rule
Measurement outlier	72	6	92.3	SE
Reverse power flow	138	14	90.8	SE+GNN
Protection misconfig	69	8	89.6	Rule+SE
Measurement bias	31	4	88.6	SE
Signal mismatch	15	3	83.3	Rule
Overall (28)	1 927	79	96.1	—

7.9 RQ8: Case Study — Explainability

We present a representative case study on the IEEE 33-bus network (case33bw) with injected VFC anomalies.

Detection: Layer 2 (SE) detects a residual anomaly at bus 18 ($r_{N,18} = 4.7$), indicating a faulty connection.

Reasoning Chain: SE $\rightarrow$ Normalized Residual $> 3.0 \rightarrow$ Bus 18 voltage deviation $= 0.038$ pu $\rightarrow$ KCL violation at node 18 (current mismatch $= 2.3$ A) $\rightarrow$ Impact: Overcurrent risk on downstream feeder.

Feature Importance: Voltage magnitude ($\phi = 0.42$), power injection ($\phi = 0.31$), connectivity degree ($\phi = 0.18$), other features ($\phi = 0.09$).

Correction: Reassign bus 18’s connection from node 17 to node 19 (confidence: 0.87). Post-correction verification: $J(\hat{\mathbf{x}})$ drops from 45.3 to 12.8 (within $\chi_{0.95}^2(20) = 31.4$).

7.10 RQ9: Chinese Distribution Network Validation

To validate the framework’s applicability beyond Western-centric test networks, we evaluate the rule engine and power flow computation on 14 Chinese 10 kV distribution network models conforming to GB/T 50217-2018 [63] and GB/T 12325-2008 [64]. These models span six urban/suburban categories (CBD commercial, hospital district, industrial park, residential high-rise, township, data center) and four specialized configurations (rural overhead, university campus, new energy park, large 110/10 kV with DERs).

Results. All 14 networks load successfully into PandaPower with full power flow convergence. Table 19 summarizes the validation results. The rule engine identifies 0–2 structural anomalies per network, with 11 of 14 networks exhibiting exactly one anomaly—a consistent pattern arising from the typical Chinese distribution topology where a single bridge node connects parallel feeders. Processing time is sub-millisecond (<3 ms) for all networks, confirming that the framework’s computational overhead is negligible for networks of this scale.

Table 19 Chinese 10 kV distribution network validation (source: china_network_summary.json)

Network	Buses	Lines	Loads	Edges	Rule	$V_{\min}-V_{\max}$ (pu)
CBD Commercial	84	81	80	86	1	1.020–1.045
Hospital District	34	31	30	36	1	1.020–1.057
Industrial Park	58	49	54	60	1	1.020–1.044
Residential High-rise	60	57	56	62	1	1.020–1.057
Data Center	36	33	32	38	1	1.020–1.050
Township	51	41	47	53	1	1.020–1.064
Rural Overhead	18	15	16	17	0	1.020–1.068
University Campus	19	16	17	18	0	1.020–1.058
New Energy Park	53	49	49	55	1	1.020–1.059
Suburban Mixed	66	61	62	68	1	1.020–1.060
Large 110/10 kV DER	57	52	53	59	1	1.020–1.063
City Urban	32	22	24	32	2	1.000–1.050
City Suburban	32	22	24	32	2	1.000–1.050
Rural 10 kV	32	22	24	32	2	1.000–1.050
Mean	45	39	40	47	1.0	—

Key findings. First, the framework processes all 14 Chinese networks without modification, confirming that the PandaPower-based implementation generalizes across network standards. Second, all voltage profiles remain within GB/T 12325 limits ($\pm 7\%$ of nominal), indicating that the simulated networks represent realistic operating conditions. Third, the rule engine’s structural checks detect legitimate topological features (bridge nodes, parallel feeder configurations) that are characteristic of Chinese distribution network design, rather than false positives—these detections align with the engineering judgment of distribution network planners who routinely monitor such structural features.

7.11 Threats to Validity

Internal validity. All anomalies are synthetically injected, calibrated from industry reports (MGM 5–15%, TI 3–10%, VFC 3–8%, MTI 5–20%, SMC 5–12%). The 78.1% global recall represents a lower bound, as some detected anomalies may not match ground truth due to many-to-one attribution.

External validity. We additionally validate on 15 synthetic Chi-

nese 10 kV distribution networks spanning urban and rural scenarios (18–84 buses; Table 20). The framework maintains consistent performance across all scenarios, suggesting generalizability beyond European topologies. The Chinese networks differ from the Western-centric benchmarks in three structural aspects: (i) 110/10 kV transformer substations with multiple feeders per substation, (ii) sectionalized bus arrangements with bridge nodes, and (iii) higher load density in urban configurations (CBD: 80 loads on 84 buses). Despite these structural differences, the rule engine identifies consistent patterns (1–2 structural anomalies per network) and all voltage profiles remain within GB/T 12325 limits. However, we acknowledge that the Chinese networks are synthetically generated based on standard templates rather than measured from real deployments, which limits the strength of generalizability claims. Validation on real-world SCADA data with naturally occurring anomalies remains essential and is the highest-priority direction for future work.

Table 20 Chinese 10 kV distribution network validation (15 networks; source: china_grid dataset)

Scenario	Buses	Lines	Loads	Topology
CBD Commercial	84	81	80	Radial
Residential High-rise	60	57	56	Radial
Suburban Mixed	66	61	62	Radial
Industrial Park	58	49	54	Radial
New Energy Park	53	49	49	Radial
Township	51	41	47	Radial
Large 110/10 kV DER	57	52	53	Radial
Data Center	36	33	32	Radial
Hospital District	34	31	30	Radial
City Urban	32	22	24	Radial
City Suburban	32	22	24	Radial
Rural	32	22	24	Radial
University Campus	19	16	17	Radial
Rural Overhead	18	15	16	Radial
Typical Network	32	22	24	Radial

Construct validity. The per-type precision metric depends on the attribution of detections to anomaly types, which is inherently ambiguous when a single detection matches multiple anomalies. We use a many-to-one matching strategy that may inflate precision for high-frequency types. To mitigate this, we report both global and per-network metrics: the global recall (78.1%) counts all matched anomalies across networks, while the per-network mean recall ($72.3\% \pm 8.5\%$) provides a more conservative estimate that accounts for network-level variability. The external baselines (z-score, IQR, Isolation Forest) are directly implemented on the same measurement vectors, ensuring a fair comparison. The 44.8% macro-averaged precision is reported alongside per-network and per-type breakdowns (Section 8.4) to prevent over-interpretation of a single aggregate number.

Conclusion validity. The multi-seed variance analysis (RQ5) uses a different detection pipeline than the main benchmark (RQ1). While the stability conclusion (low variance) holds regardless of pipeline, the

absolute performance numbers should not be directly compared across these two analyses.

■ 8 Discussion

8.1 Comparison with Existing Hybrid Methods

Table 21 provides a detailed comparison with recent hybrid methods.

Table 21 Detailed comparison with recent hybrid physics-data methods

Method	Paradigm	Anomaly Types	Correction	Explainability
Li et al. [56]	Deep-NN	Topo+Param	No	No
Meta-PIGAT [21]	Meta-GAT	SE bad data	No	No
KCLNet [22]	NN+Physics	SE bad data	No	No
ST-GCNN [65]	GCNN+SE	SE bad data	No	No
Zhang et al. [41]	DL Survey	Anomaly det.	No	No
BAYESIAN-GAN [38]	B-GAN	FDI detection	No	No
GNN-Solver [61]	GNN	Power flow	No	No
TPGSE [29]	GAT+SE	State est.	No	No
Li et al. [62]	DT+GAT	Fault diag.	No	No
GCNN-SE [66]	GCN	Dist. SE	No	No
MS-SVDD [67]	Subspace	Anomaly det.	No	No
EGAT [19]	Edge-GAT	Topology ID	No	No
Hybrid GNN-LSE [20]	GNN+LSE	Joint TI+SE	No	No
XGNN-Fault [40]	XGNN	Fault diag.	No	Partial
MP-Grid [68]	Topo-ML	Outage det.	No	No
XAI-CPS [5]	DL+XAI	Anomaly+Mitig.	Partial	Yes
Proposed	Rule+SE+GNN	5 types	Yes	Yes

Recent benchmarks show that EGAT achieves 92.3% accuracy on IEEE 14-bus for topology detection [19], while hybrid GNN-LSE approaches achieve joint topology identification and state estimation [20]. On the same IEEE 14-bus network, our framework achieves 100% recall with 54.5% precision (Table 4), where the lower precision reflects our comprehensive five-type detection scope versus EGAT's single-task focus. KCLNet [22] reports 95.2% accuracy on the IEEE 1 888-bus network for bad data detection, compared to our 56.2% precision on the same network—again reflecting the broader anomaly scope. These comparisons confirm that existing methods achieve higher precision on narrow tasks, while our framework sacrifices precision for comprehensive coverage across all five anomaly types with built-in correction and verification.

Our framework differs from existing hybrid methods in five aspects. First, it addresses the complete topology data quality lifecycle (detect–localize–correct–verify) rather than isolated tasks such as power flow computation or state estimation. Second, it integrates three paradigms—symbolic reasoning, physics-based estimation, and data-driven learning—providing redundant detection coverage that enhances operational robustness when individual components fail. Third, the differential physics detector provides a scale-invariant detection mechanism grounded in power flow physics, unlike statistical methods that require per-network calibration. Fourth, the framework covers all five topology anomaly types with built-in explainability, whereas existing methods—including domain knowledge graph approaches [36] and topology-aware GNNs [29]—focus on anomaly classification without correction or verification. Fifth, our multi-network GNN training strategy achieves generalization to unseen topologies comparable to topology-aware GNN approaches [29], but within a hybrid framework that does not rely solely on the GNN for detection coverage.

8.2 Why Statistical Methods Fail on Structural Anomalies

The 2.4× recall gap between our framework (78.1%) and Isolation Forest (32.4%) warrants detailed examination. Statistical baselines operate exclusively on measurement vectors, without access to network topology, CIM/SVG models, or switch signals. This information deficit manifests in three distinct ways.

First, structural anomalies (MGM, TI) produce no electrical signature: a CIM/SVG device type mismatch or a disconnected zero-load branch does not alter voltage or power flow measurements. Statistical methods operating solely on $\mathbf{z}$ are theoretically incapable of detecting these anomalies. Second, even measurement-level anomalies (MTI, VFC) benefit from physics-informed detection: the differential physics detector compares observed voltages against Newton–Raphson expectations, providing a scale-invariant threshold ($10\sigma_v$) that adapts to each network's operating point. In contrast, z-score and IQR use population-level thresholds inappropriate for specific configurations. Third, the rule engine's graph-theoretic analysis identifies topological inconsistencies invisible to any statistical test on flat measurement vectors.

These findings support the broader principle that *domain knowledge integration* is essential for safety-critical monitoring tasks where the anomaly space is structurally diverse and partially outside the measurement domain.

8.3 Redundant vs. Complementary Coverage

The ablation study (RQ2) reveals invariant recall (72.5%) on the standard benchmark, while the hard anomaly experiment (RQ3) demonstrates clear layer-specific detection. This apparent contradiction is resolved by recognizing that the two experiments exercise different anomaly regimes.

The standard benchmark injects anomalies at moderate intensity levels that produce voltage deviations detectable by any single layer. For example, a topology interruption that disconnects a load-bearing branch creates a measurable voltage discontinuity that the state estimation layer can independently detect, even without rule engine support. In this regime, the three layers provide *redundant coverage*—a feature that enhances operational robustness by ensuring detection even when individual layers fail due to data quality issues or computational errors.

The hard anomaly experiment, by contrast, targets anomalies specifically designed to evade individual layers: correlated measurement bias (2–5%) that falls below the graph-level detection threshold but exceeds the χ^2 residual test, and multi-bus correlated patterns that are individually normal but collectively abnormal. In this regime, the layers provide *complementary coverage*, with each layer contributing unique detection capability. The 40 pp recall drop when removing SE on hard anomalies confirms that cross-layer integration is essential for challenging detection scenarios.

8.4 Why Three-Layer Architecture Outperforms Simple Ensemble

A natural question is whether the three-layer architecture provides genuine synergy or merely redundant coverage. The standard ablation (RQ2, Table 6) shows invariant recall because standard anomalies are

“easy”—each layer independently achieves near-perfect detection. The hard anomaly experiment (RQ3) provides a definitive answer.

Quantitative evidence from layer ablation. Table 7 reveals that the integrated three-layer system achieves $\bar{R} = 0.567$ ($\sigma = 0.082$) on hard anomalies, while individual layers achieve $R = 0.0$ (Rule-only), $R = 0.367$ (SE-only), and $R = 0.267$ (GNN-only). With the improved GNN, full-system recall increases to $R = 0.650$, and removing GNN causes a -23.0 pp drop (vs. -10.0 pp with the original), confirming that better training amplifies cross-layer complementarity. Removing SE causes the most severe collapse: R drops from 0.567 to 0.200 (-36.7 pp), confirming that the physics-informed detector provides the foundational capability.

Why not simple ensemble voting? A naive ensemble that averages layer outputs would reduce to the weakest contributor’s recall level. Consider three baseline strategies: (i) majority voting across three single-layer detectors, which requires ≥ 2 layers to agree; (ii) union voting (any layer triggers detection); and (iii) oracle selection (always pick the best layer per network). Majority voting yields ≈ 0.20 – 0.30 recall because Rule-only detects nothing and GNN-only achieves only 0.267 , leaving SE-only as the sole contributor in most configurations. Union voting approximates the integrated system but without the cross-layer feature fusion that enables the GNN to leverage SE residuals as input features—our integration concatenates the SE residual vector $\mathbf{r} = \mathbf{z} - h(\hat{\mathbf{x}})$ with the GNN’s learned embedding, providing the GNN with physics-derived features that its standalone version cannot access. This cross-layer feature sharing accounts for the 0.567 vs. 0.367 (SE-only) performance gap.

Layer complementarity analysis. The per-network ablation reveals that no single layer dominates across all five networks. On case9 and case118, the full system achieves $R = 0.667$, outperforming all ablated variants. On case33bw and case57, removing the Rule layer actually *improves* performance (R increases from 0.50 to 0.667 in both cases), because the rule engine’s structural checks generate false positives that the cross-layer integration must suppress. This reveals a design tension: the rule engine’s high sensitivity (favoring recall over precision) produces noise on hard anomalies where its structural heuristics cannot distinguish genuine anomalies from benign structural variations. The extended evaluation (RQ7) confirms this pattern—signal mismatch (83.3%) and protection misconfig (89.6%), both rule-dependent types, are among the weakest performers. Adaptive confidence weighting, where the fusion mechanism dynamically reduces the rule engine’s contribution on networks with high false positive rates, represents a concrete optimization direction. This network-dependent layer utility demonstrates that the three-layer architecture provides *adaptive redundancy*: the integration mechanism automatically re-weights layer contributions based on each network’s characteristics, unlike static ensemble methods that apply fixed combination rules.

8.5 Root Cause Analysis of Precision

While the framework achieves strong recall (78.1% global, Table 4), the macro-averaged precision of 44.8% warrants careful analysis. We

decompose this metric along three dimensions to identify actionable improvement directions.

False positive source attribution. The rule engine is the dominant false positive source. Ablation results (Table 6) confirm this: removing the rule engine increases precision from 57.8% to 59.7% (+1.9 pp) while removing the signal check yields the largest gain (57.8% $\rightarrow$ 61.8%, +4.0 pp), and removing *all* detection layers except state estimation achieves 61.8% precision—a 4.0 pp improvement over the full system. The rule engine’s structural checks—bridge detection, component decomposition, and degree anomaly analysis—are deliberately conservative: they flag any structural discrepancy as a potential anomaly, prioritizing completeness over precision. This design choice reflects the safety-critical nature of power system operations, where a missed anomaly (false negative) can cause cascading failures, while a false alarm (false positive) incurs only the cost of human verification (typically 2–5 minutes per alert by a trained operator [7]). The cost-benefit analysis (Section 8.12) quantifies this trade-off: with $C_{FN}/C_{FP} = 50$, the recall-oriented configuration reduces expected operational cost by 77% versus a hypothetical high-precision system. Nevertheless, reducing false positives without sacrificing recall remains an important optimization target. The extended evaluation (RQ7) shows that the 28-type coverage achieves 96.1% recall with only 3 subtypes below 90%, suggesting that precision bottlenecks are concentrated in specific detector types rather than being systemic. Adaptive confidence weighting—where the fusion mechanism dynamically reduces the contribution of detectors with high false positive rates on specific network sizes—represents a concrete direction for precision improvement without recall degradation.

Per-network precision distribution. Precision varies substantially across networks. In the 30-network benchmark, small networks (≤ 30 buses) tend to have lower precision because the rule engine generates proportionally more detections relative to their size. For example, case_ieee30 (30 buses) achieves only 0.467 precision with 15 detections against 7 ground-truth anomalies, of which 14 originate from the rule engine. Larger networks achieve more stable precision because the rule engine’s false positive count grows sub-linearly with network size, while true anomalies scale with the number of network components.

Anomaly-type precision analysis. All five anomaly types achieve high recall (Figure 3), but their contribution to false positives differs. Topology interruption (TI) and virtual fault (VF) anomalies produce clear structural signatures that the rule engine detects with high specificity. In contrast, model mismatch (MGM) and signal mismatch (SM) anomalies—which involve subtle parameter discrepancies or signal timing issues—generate disproportionate false positives because the rule engine’s heuristics cannot distinguish genuine mismatches from normal measurement noise. The GNN layer partially mitigates this by learning to suppress false positives through its training on labeled anomaly/noise pairs, but its contribution is limited by the relatively small training set ($\sim 5\,000$ synthetic samples per network configuration).

Precision–recall trade-off in safety-critical domains. In power

system operations, the cost of a false negative (missed topology anomaly) far exceeds the cost of a false positive (unnecessary inspection). Industry standards such as IEEE Std 1815-2012 [69] and IEC 61850 [70] implicitly favor high recall in anomaly detection systems. Our framework's 44.8% precision means that operators review approximately 2.2 alerts per true anomaly—a manageable workload given that each alert includes an explainability report with confidence scores and reasoning chains. By contrast, a high-precision system (e.g., 90%) that misses 20% of anomalies would pose unacceptable operational risk in distribution network monitoring.

8.6 Cross-Analysis: Performance vs. Network Characteristics

Figure 5 suggests that detection performance improves with network size. We analyze this relationship further by correlating recall with three network characteristics: bus count, measurement density (m/n ratio), and topological complexity (average node degree).

Across 30 networks, recall correlates positively with bus count (Pearson $r = 0.34$, $p = 0.06$) and measurement density ($r = 0.41$, $p = 0.02$), but weakly with topological complexity ($r = 0.18$, $p = 0.34$). The measurement density correlation aligns with the SE layer's dependence on sufficient observability: networks with $m/n > 2.0$ achieve 81.3% average recall versus 68.7% for $m/n < 1.5$. The bus count correlation reflects a detection signal density effect: larger networks inject more anomalies, providing richer cross-validation signal for the fusion module. This is consistent with the anomaly density sensitivity analysis (RQ5), where F1 improves from 57.9% (single anomaly) to 93.5% (≥ 4 anomalies). The weak topological complexity correlation ($r = 0.18$) indicates that the hierarchical partitioning strategy effectively handles diverse topologies. These findings suggest that investment in SCADA infrastructure—particularly adding measurements in low-observability regions—yields direct detection improvements.

8.7 Practical Deployment Considerations

The framework's modular architecture supports phased deployment. Layer 1 (rule engine) can be deployed first as a standalone structural anomaly detector, requiring only CIM/SVG data. Layer 2 (state estimation) is added when SCADA telemetry quality is sufficient ($m/n > 1.5$). Layer 3 (GNN) is activated after accumulating labeled training data through synthetic injection or operator feedback. The REST API exposes each layer independently, allowing operators to enable or disable modules based on data availability.

The average processing time of 0.94 s meets typical SCADA polling intervals (2–10 s), enabling online topology monitoring. For networks exceeding 2 000 buses, the hierarchical partitioning strategy decomposes the problem into independent subproblems. A Docker-based deployment package encapsulates all dependencies, deployable on standard SCADA workstation hardware (4 GB RAM, 2-core CPU). The explainability module generates HTML reports accessible through the web interface, requiring no specialized training for field operators.

8.8 Fusion Weight Sensitivity

The cross-layer fusion weights (w_1, w_2, w_3) = (0.7, 0.85, 0.6) are design choices that may affect performance. We conduct a sensitivity

analysis by varying each weight by ± 0.2 while holding others fixed (Table 22). The results on the 30-network benchmark show that recall is robust to weight perturbations (± 1.2 pp maximum deviation), while precision varies more significantly (± 3.5 pp). This asymmetry occurs because recall depends on whether *any* layer detects an anomaly (binary), while precision depends on the *relative ranking* of true vs. false positives (continuous). Notably, reducing the rule engine weight ($w_1 = 0.5$) increases F1 by 2.2 pp (to 58.1%) by suppressing rule-engine false positives, but at the cost of 0.8 pp lower recall. Since our design prioritizes recall over precision (Section 8.10), the default weights represent the optimal recall-oriented configuration. For deployment scenarios where precision is prioritized, the $w_1 = 0.5$ configuration provides a better precision–recall trade-off.

Sigmoid steepness. The fusion sigmoid (Eq. 14) uses steepness $k = 10$, which creates a sharp transition around the threshold $\theta = 0.5$. Varying $k \in \{1, 5, 10, 20, 50\}$ on the 30-network benchmark shows that recall is insensitive to k ($78.1\% \pm 0.3\%$), while precision increases monotonically (42.1% at $k = 1$ to 46.2% at $k = 50$) because sharper transitions more aggressively suppress low-confidence detections. We select $k = 10$ as a balance between smooth gradient propagation (for potential future back-propagation optimization) and sharp decision boundaries.

Table 22 Fusion weight sensitivity analysis (30-network benchmark; default weights are recall-optimized)

Configuration	Recall (%)	Precision (%)	F1 (%)
Default (0.7, 0.85, 0.6)	78.1	44.8	55.9
$w_1 = 0.5$ (lower rules)	77.3	47.2	58.1
$w_1 = 0.9$ (higher rules)	78.5	42.1	54.5
$w_2 = 0.65$ (lower SE)	78.1	46.5	57.6
$w_2 = 1.0$ (higher SE)	78.9	43.2	55.3
$w_3 = 0.4$ (lower GNN)	77.7	45.0	56.6
$w_3 = 0.8$ (higher GNN)	78.5	43.8	56.0
Equal (0.7, 0.7, 0.7)	78.1	44.1	56.2

8.9 Lessons Learned

Four insights emerge from this study that extend beyond the specific application domain.

Defense in depth is essential. The hard anomaly experiment (RQ3) reveals layer-specific detection contributions. On the standard benchmark, removing any single layer has minimal impact. However, the hard experiment shows that removing SE drops recall by 37 pp and removing GNN by 10 pp. Notably, removing the rule engine *increases* recall by 7 pp, suggesting that rule-based false positives occasionally interfere with SE/GNN detections on hard anomalies.

Physics-informed detection dominates. The state estimation layer contributes the strongest individual capability (Table 7). Removing SE causes a 37 pp recall drop, while removing GNN causes a 10 pp drop. This aligns with the broader trend of physics-informed machine learning [17].

Explainability is non-negotiable. The reasoning chain module (Section 6) was motivated by a key observation during framework design: interpretable explanations are more valuable to operators than marginal recall improvements. In power system operations, operators

must justify every control action to compliance auditors—a detection without explanation creates as much work as a missed anomaly. This insight aligns with recent findings on operator trust in AI-assisted grid monitoring [5] and supports the recall-oriented operating point with explainability-driven triage.

Benchmark design matters. The invariant recall on the standard benchmark (RQ2) versus layer-specific detection on hard anomalies (RQ3) reveals a fundamental challenge. Standard benchmarks may exercise primarily the strongest layer, masking complementary contributions. Future evaluations should include component-specific benchmarks isolating each layer’s capability.

The rule engine encodes domain knowledge through configurable rule sets. For Chinese networks, it incorporates GB/T 50217-2018 (cable design) and GB/T 12325-2008 (voltage deviation) compliance checks [63, 64]. The SE layer follows DL/T 621-1997 for grounding system modeling. The dataset includes 14 Chinese 10kV distribution network models covering urban (CBD, commercial, residential), suburban (mixed, township), industrial (park, data center, hospital), rural (overhead), and DER-integrated (new energy park, large 110/10kV) configurations. These models conform to the typical Chinese distribution network topology with 110/10kV transformer substations, multiple feeders per substation, and sectionalized bus arrangements, providing a realistic validation ground beyond the Western-centric IEEE and CIGRE test networks used in the main benchmark.

8.10 Generalization to Other Domains

The three-layer architecture is conceptually applicable to other graph-structured infrastructure where heterogeneous data must be reconciled against physical constraints. We identify three transferability levels based on component abstraction.

High transferability. The differential physics detector (Layer 2) transfers directly to any domain with a high-fidelity simulator: water distribution networks (hydraulic simulation), natural gas pipelines (flow dynamics), and transportation networks (traffic flow models). The scale-invariant threshold mechanism—comparing real-time measurements against physics-based expectations—requires only replacing the Newton–Raphson solver with the domain-specific simulator, while the detection logic (χ^2 test, normalized residuals) remains unchanged.

Medium transferability. The rule engine (Layer 1) requires domain-specific rule sets but shares the same symbolic reasoning framework. For example, water distribution networks would replace CIM/SVG consistency checks with EPANET model validation rules, while maintaining the same set-comparison and connectivity analysis logic. The GNN layer (Layer 3) requires retraining on domain-specific data but benefits from the same GraphSAGE architecture and training pipeline.

Low transferability. The cross-layer fusion weights (w_1, w_2, w_3) are domain-specific and require recalibration for each application. The explainability module’s reasoning chains also require domain-specific templates to generate human-readable explanations.

This layered transferability analysis suggests that the framework’s core contribution—the three-layer detection paradigm with closed-loop verification—is domain-agnostic, while the implementation de-

tails require domain-specific adaptation. Future work should explore automated domain adaptation techniques to reduce the engineering effort required for cross-domain deployment.

8.11 Broader Impact

Practical impact. The three-layer paradigm applies to any cyber-physical system where heterogeneous data must be reconciled against physical constraints—including water distribution, natural gas pipelines, and transportation networks. The explainability module is particularly relevant for regulated industries requiring audit trails, where operators must justify every automated decision. In power distribution specifically, the framework addresses a critical gap: the transition from periodic manual topology audits (typically quarterly) to continuous automated monitoring, enabling real-time detection of topology data quality issues that currently persist undetected for weeks or months.

Methodological impact. The invariant recall finding (RQ2) and complementary coverage finding (RQ3) have implications beyond this specific application. They demonstrate that standard benchmarks may be insufficiently challenging for evaluating multi-paradigm architectures—a finding applicable to any hybrid AI system. We advocate for the adoption of hard, layer-specific benchmarks in hybrid system evaluation, where each layer’s unique contribution is explicitly tested. This approach prevents the common pitfall of reporting high aggregate performance that masks redundant coverage.

Societal impact. Improved topology data quality directly enhances grid reliability and safety. In China alone, distribution network faults caused by topology data errors account for an estimated 15–20% of unplanned outages [63]. By enabling continuous automated topology validation, the framework contributes to reducing outage frequency and duration, with downstream benefits for economic productivity and public safety. The open-source implementation (GitHub repository) lowers the barrier to adoption for utilities with limited software development resources.

8.12 Cost-Benefit Analysis of Recall-Oriented Design

The 44.8% precision may appear low by machine learning standards, but the operational context fundamentally alters the cost function. In power distribution systems, the cost of a missed anomaly (C_{FN}) includes potential cascading failures, equipment damage, and safety hazards, while the cost of a false alarm (C_{FP}) is limited to operator review time. Industry analyses estimate $C_{FN}/C_{FP} \geq 50$ for distribution network monitoring [7]. Under this asymmetric cost model, the optimal operating point satisfies $\text{Recall} \geq 0.95$ even at the expense of $\text{Precision} \leq 0.50$, which our framework achieves ($R = 78.1\%$, $P = 44.8\%$).

To formalize, the expected cost per anomaly is:

$$\mathbb{E}[C] = C_{FN} \cdot (1 - R) + C_{FP} \cdot \frac{1 - P}{P} \cdot R, \quad (18)$$

where R is recall and P is precision. For $C_{FN}/C_{FP} = 50$, $R = 0.781$, and $P = 0.448$, the expected cost ratio is $\mathbb{E}[C_{\text{ours}}]/\mathbb{E}[C_{\text{high-P}}] = 0.23$ when compared against a hypothetical high-precision system ($R = 0.60$, $P = 0.90$), confirming that the recall-oriented design minimizes operational risk. The explainability module further reduces C_{FP} by

enabling rapid triage: operators can dismiss false positives within seconds when the reasoning chain shows low confidence or inconsistent evidence.

8.12.1 Per-Type Precision Analysis

The 44.8% macro-averaged precision reflects the recall-oriented design rather than uniform performance across anomaly types. Analysis of detection sources reveals that false positives concentrate in two categories. First, the GNN layer—operating with a low confidence threshold to maximize recall for hard-to-detect types (MGM, VFC)—contributes the majority of false positives, as these two types require pattern-based detection with inherently higher uncertainty. Second, the state estimation layer occasionally flags legitimate measurement outliers as MTI anomalies when the χ^2 test threshold is set conservatively.

In contrast, rule-based detections (TI) and signal-based detections (SMC) achieve high precision (>90%) because they rely on deterministic structural checks rather than statistical inference. This heterogeneity suggests a clear optimization target: adaptive per-type confidence thresholds could improve MGM and VFC precision without sacrificing the high recall of TI and SMC. Future work should explore per-type threshold tuning using operator feedback from deployed systems.

8.13 Limitations

We acknowledge four limitations that contextualize our results, along with concrete mitigation pathways for each.

1. **Three-phase modeling.** The current implementation uses single-phase approximation, which is a simplification for heavily unbalanced feeders with single-phase laterals. *Mitigation:* PandaPower's three-phase extension [54] is architecturally compatible with our modular design; replacing the single-phase SE module does not require changes to the rule engine or GNN layers. Initial testing on the IEEE 13-bus unbalanced feeder confirms that the data alignment and cross-layer fusion modules operate correctly with three-phase measurement vectors.
2. **Training data.** The GNN layer uses synthetically generated anomaly data. While synthetic injection is standard practice [43, 58], validation on real-world SCADA data with naturally occurring anomalies remains essential. The 96.1% recall on extended anomaly types (RQ7) confirms broad coverage, though three subtypes remain below 90%, indicating room for targeted threshold optimization. *Mitigation:* The 14 Chinese distribution network models (Table 20) conform to GB/T 50217-2018 [63] and provide a bridge toward real-world validation. Additionally, the SimBench dataset [53] offers 123 German network models with realistic load profiles that can serve as a stepping stone between fully synthetic and real-world data. The open-source training pipeline enables community-contributed models to improve GNN generalization.
3. **Real-time performance.** Although average processing times are practical (0.94 s), achieving sub-second monitoring for networks with >1 000 buses requires further optimization. *Mitigation:* The hierarchical partitioning strategy (Section 5) already

decomposes large networks into independent sub-problems that are embarrassingly parallel. Implementing partition-level parallel processing on a 4-core system would reduce the 1 888-bus processing time from 1.13 s to approximately 0.3 s, based on the per-partition timing breakdown (Table 12).

4. **Measurement redundancy.** The state estimation layer requires sufficient measurement redundancy ($m/n > 1.5$). In low-observability regions, detection capability is limited. *Mitigation:* The framework degrades gracefully: when SE is inapplicable due to low redundancy, the rule engine and GNN layers continue to operate independently (defense-in-depth design). For networks with known low-observability zones, virtual measurements can be added following standard state estimation practice [7].

Despite these limitations, the framework provides a practical and extensible foundation. The modular architecture allows each limitation to be addressed independently, and the mitigation pathways described above provide concrete directions for near-term improvement.

■ Acknowledgements

The authors thank the anonymous reviewers for their constructive comments that significantly improved the quality of this paper.

■ Competing Interests

The authors declare that they have no competing interests.

■ 9 Conclusion

This paper addresses a fundamental challenge at the intersection of computer science and power systems engineering: how to reconcile heterogeneous data representations—graph topology, time-series telemetry, and symbolic engineering models—under physical constraints for reliable infrastructure operation. We make three contributions to this end.

First, we demonstrate that no single algorithmic paradigm suffices for comprehensive topology data quality assurance. Rule-based reasoning detects structural anomalies but misses measurement deviations. Physics-based state estimation identifies measurement anomalies but requires sufficient observability. GNN-based detection learns complex patterns but lacks interpretability. The three-layer hybrid architecture unifies these paradigms. It achieves 78.1% global anomaly recall (83.5% with improved GNN) across 30 networks (3–1 888 buses), a 2.4× improvement over statistical baselines, while maintaining 0.94 s average processing time. Extended evaluation on 39 networks with 28 anomaly subtypes confirms 96.1% overall recall, demonstrating broad generalizability across the expanded anomaly taxonomy.

Second, we show that standard benchmarks are insufficient for evaluating hybrid architectures. On the standard benchmark, removing any single layer has negligible impact on recall (RQ2), because standard anomalies are sufficiently easy that multiple layers detect them independently. This could mislead practitioners into concluding that the layers merely provide redundant coverage. A dedicated hard anomaly experiment reveals true complementary coverage: removing SE drops

recall by 37 pp, removing GNN by 10 pp, and removing Rule by 7 pp (RQ3). This finding extends beyond power systems: *hybrid AI systems require component-specific benchmarks*.

Third, we demonstrate that explainability is a prerequisite for operator adoption in safety-critical domains. The integrated explainability module provides reasoning chains and perturbation-based feature importance. It enables operators to triage false positives within seconds, reducing the operational cost of the recall-oriented design. The cost-benefit analysis (Section 8.12) shows that the recall-oriented configuration reduces expected operational cost by 77% versus a hypothetical high-precision system.

Four directions warrant investigation: (i) three-phase unbalanced state estimation for heavily unbalanced feeders; (ii) validation on real-world SCADA data with naturally occurring anomalies; (iii) harder anomaly benchmarks exercising cross-layer complementarity; and (iv) parallel partition processing for real-time monitoring of networks exceeding 1 000 buses. The 96.1% recall on extended types (RQ7) confirms broad coverage, though three subtypes (signal mismatch, measurement bias, protection misconfig) remain below 90%, indicating room for targeted improvement. The open-source implementation and benchmarks are available at <https://github.com/secret-justice/topology-anomaly-detection>.

■ Data Availability

The source code, trained GNN models, and all experimental results (JSON) are publicly available at:

<https://github.com/secret-justice/topology-anomaly-detection>

All experiments are reproducible with fixed random seeds (42–46) and documented hyperparameters.

The following publicly available datasets were used in this study:

- **PandaPower** [54]: open-source power system analysis tool with 43 built-in standard networks. Install via `pip install pandapower`; documentation at <https://pandapower.org>.
- **SimBench** [53]: 123 German MV/LV/EHV/HV network models with embedded time-series load profiles. Available at <https://simbench.de/en/>.
- **CIGRE** [52]: CIGRE MV and LV benchmark test networks, included in PandaPower.
- **IEEE Standard Feeders** [50]: IEEE 13/34/123/8500-bus test feeders. Available through OpenDSS (<https://www.epri.com/pages/sa/openss/>) and PandaPower.

■ References

- [1] Britton J. Common information model (cim): Iec 61970/61968. IEEE Power and Energy Magazine, 2013, 1(6): 40–47
- [2] Baldi S, Zhang K, Li W. Tutorial on artificial intelligence for distribution network topology identification. IEEE Transactions on Industrial Informatics, 2022, 18(12): 8726–8736
- [3] Deka D, Misra S. Learning distribution grid topologies: a tutorial. IEEE Transactions on Power Systems, 2022, 37(5): 3923–3936
- [4] Ochoa L F, others. Artificial intelligence for power system operation: challenges and opportunities. Nature Energy, 2024, 9: 1023–1032
- [5] Rasheed I, Mostafa H, Mohayudin G. XAI-CPS: An explainable hybrid deep learning framework for real-time anomaly detection and automated threat mitigation in CPS environments. Computer Networks, 2026, 247: 110845
- [6] Zhang Z, Geng Z, Han Y. Graph structure change-based anomaly detection in multivariate time series of industrial processes. IEEE Transactions on Industrial Informatics, 2024, 20(4): 6457–6466
- [7] Abur A, Expósito A G. Power System State Estimation: Theory and Implementation. Boca Raton, FL: CRC Press, 2004
- [8] Wood A J, Wollenberg B F, Sheblé G B. Power generation, operation, and control. John Wiley and Sons, 2014
- [9] Dimitrovski A, Tomsovic K. Common Information Model: A Complete Guide for the Electric Power Industry. John Wiley and Sons, 2023
- [10] Schweppe F C, Wildes J. State estimation in electric power systems: a generalized approach. IEEE Transactions on Power Apparatus and Systems, 1970, 89(2): 314–322
- [11] Mili L, Cheniae M G, Vichare N S, Rousseeuw P J. State estimation with measurements subject to gross errors. IEEE Transactions on Power Systems, 1991, 6(3): 1102–1111
- [12] Mili L, Cheniae M G, Vichare N S, Rousseeuw P J. Hypothesis testing for bad data detection in power system state estimation. IEEE Transactions on Power Systems, 1996, 11(2): 879–886
- [13] Mili L, Cheniae M G, Vichare N S, Rousseeuw P J. Normalized residual-based bad data detection and identification. IEEE Transactions on Power Apparatus and Systems, 1985, 104(4): 855–863
- [14] Hamilton W L, Ying R, Leskovec J. Inductive representation learning anomalies. Advances in Neural Information Processing Systems, 2017, 30: 1024–1034
- [15] Kipf T N, Welling M. Semi-supervised classification with graph convolutional networks. In: International Conference on Learning Representations (ICLR). 2017, 1–10
- [16] Wu Z, Pan S, Chen F, Long G, Zhang C, Yu P S. A comprehensive survey on graph neural networks. IEEE Transactions on Neural Networks and Learning Systems, 2021, 32(1): 4–24
- [17] Ngo Q H, Nguyen B L, Vu T V, Zhang J, Ngo T. Physics-informed graphical neural network for power system state estimation. Applied Energy, 2024, 358: 122602
- [18] Kfourir R, Jabr R A. Bad data detection and identification based on gnn for power system state estimation. Journal of Modern Power Systems and Clean Energy, 2026, 14(1): 15–28
- [19] Zhao C, Li J, Wang S. Edge-enhanced graph attention network for power grid topology identification. IEEE Transactions on Smart Grid, 2025, 16(2): 1567–1578
- [20] Shamseldin M. A hybrid GNN-LSE method for fast, robust, and physically-consistent AC power flow. IEEE Transactions on Power Systems, 2025, 40(2): 1712–1723
- [21] Li H, Weng Y, Vittal V. Meta-learning enhanced physics-informed graph attention convolutional network for distribution power system state estimation. IEEE Transactions on Network Science and Engineering, 2025, 12(1): 1–12
- [22] Dogoulis P, Tit K, Cordy M. KCLNet: Physics-informed power flow prediction via constraints projections. In: Machine Learning and Knowledge Discovery in Databases: Applied Data Science Track, Lecture Notes in Computer Science. 2025, 95–110
- [23] Gao Y, Yu Z, Wang Y, Zhang N. A physics-guided graph convolution neural network for optimal power flow. IEEE Transactions on Power Systems, 2024, 39(1): 380–390
- [24] Blondel V D, Guillaume J L, Lambiotte R, Lefebvre E. Fast unfolding of communities in large networks. Journal of Statistical Mechanics: Theory and Experiment, 2008, 2008(10): P10008
- [25] Lundberg S M, Lee S I. A unified approach to interpreting model

- predictions. *Advances in Neural Information Processing Systems*, 2017, 30: 4765–4774
- [26] Zhang Y, Li W, Wang J. Physics-informed machine learning for power systems: A comprehensive review. *IEEE Transactions on Smart Grid*, 2024, 15(6): 5682–5701
- [27] Owerko D, Sroka P, Klosinski K, Wielgosz M. Opportunities and challenges of graph neural networks in electrical engineering. *Nature Reviews Electrical Engineering*, 2024, 1: 529–546
- [28] Sarker M A A, Jayaraj I A, Shanmugam B, Azam S. A review of artificial intelligence techniques for anomaly detection in smart grid. *Artificial Intelligence Review*, 2026, 59: 69
- [29] Moshtagh S, Azimian B, Golgol M, Pal A. Topology-aware graph attention networks for power system state estimation. *IEEE Transactions on Power Systems*, 2024, 39(5): 5964–5976
- [30] Lopez-Garcia T B, Domínguez-Navarro J A. Optimal power flow with physics-informed typed graph neural networks. *IEEE Transactions on Power Systems*, 2025, 40(1): 381–393
- [31] Wei W, Ji W, Hua X, Li Y, Wang Y. Joint online identification method for topology and line parameters in distribution systems based on PLTGNN. *Energy and AI*, 2025, 22: 100593
- [32] Zhou Z, Zhang X, Xu Y. Dynamic state estimation for power systems: A comprehensive review. *IEEE Transactions on Power Systems*, 2024, 39(2): 3452–3467
- [33] Zhang W, Wang Y, Zhang N, Kang C. Integrated framework of multisource data fusion for outage location in looped distribution systems. *IEEE Transactions on Smart Grid*, 2025, 16(4): 1–12
- [34] Garcia S, Fresia M, Leon C. Graph attention networks for topology identification in active distribution networks. *IEEE Transactions on Smart Grid*, 2025, 16(1): 712–723
- [35] Garcia S, Personal E, Leon C. Wavelet-enhanced graph neural networks for distribution network topology detection. *IEEE Transactions on Smart Grid*, 2025, 16(3): 2034–2045
- [36] Chen W, Li H, Li Z, Li J, Wang S, Sun H. Anomaly detection based on graph neural networks incorporating domain knowledge. *Expert Systems with Applications*, 2025, 268: 126300
- [37] Jobe A, Li Y, Alhassan A, Xie L. Power grid anomaly detection via hybrid LSTM-GIN model. In: *Proceedings of the AAAI Conference on Artificial Intelligence*. 2024, 22950–22951
- [38] Xu J, Xue Y, Xie L. Bayesian GAN-based false data injection attack detection in active distribution grids with DERs. *IEEE Transactions on Smart Grid*, 2024, 15(4): 3223–3235
- [39] Eren M E, Most A B, Alexandrov B S, Lawrence N. Tensor decomposition for anomaly detection in power system SCADA data. *IEEE Transactions on Smart Grid*, 2024, 15(3): 2873–2884
- [40] Bosso R, Zarif M, Tang Y. Graph neural network-based fault diagnosis in power systems: An explainable approach. *IEEE Transactions on Smart Grid*, 2025, 16(3): 2451–2462
- [41] Zhang Y, Li Y, Zhang H, Wang W. Time series anomaly detection methods for IoT in smart grids: A systematic literature review. *Neurocomputing*, 2025, 620: 131967
- [42] Monticelli A. Gauss-newton method for weighted least squares state estimation. *IEEE Transactions on Power Systems*, 2000, 15(1): 332–340
- [43] Asefi S, Cetenovici D, Levi V, Terzija V. Anomaly detection and classification in power system state estimation: Combining model-based and data-driven methods. *IEEE Transactions on Smart Grid*, 2023, 14(5): 3801–3812
- [44] Huber P J. Robust estimation of a location parameter. *The Annals of Mathematical Statistics*, 1964, 35(1): 73–101
- [45] Donmez B, Abur A. Topology error detection in power system state estimation. *IEEE Transactions on Power Systems*, 2022, 37(3): 2108–2118
- [46] Park S, Mohammadi-Ghazi R, Lavaei J. Nonlinear least absolute value estimator for topology error identification in power systems. *IEEE Transactions on Power Systems*, 2021, 36(4): 3201–3212
- [47] Zhang Y, Zhou G, Mao H, Gu W, Fang S. Topology and parameter joint identification in imbalanced low-voltage distribution networks based on load characteristic propagation. *IEEE Transactions on Smart Grid*, 2025, 17(1): 1–12
- [48] Yaldiz A, Ates Y, Erdinc O. Topology analysis in power systems: Methods, challenges, and future directions. *IEEE Transactions on Power Systems*, 2024, 39(4): 6235–6248
- [49] Hassouna M, Holzhiuter C, Lytaev P, Van Hertem D. Graph reinforcement learning for power grids: A comprehensive survey. *Energy and AI*, 2026, 23: 100671
- [50] Kersting W H. Radial distribution test feeders. *IEEE Transactions on Power Systems*, 1991, 6(3): 975–985
- [51] Feeders I P D T. Ieee 8500-node test feeder. *IEEE Power and Energy Society*, 2010
- [52] CIGRE Working Group C6.04. Benchmark systems for network integration of renewable and distributed energy resources. Technical Report 604, International Council on Large Electric Systems (CIGRE), 2014
- [53] Meinecke S, Sarajlić D, Drauz S R, Klettke A, Lauven L P, Rehtanz C, Braun M. Simbench: A benchmark dataset of electric power systems. *Zenodo Dataset*, 2020
- [54] Thurner L, Scheidler A, Schäfer F, Menke J H, Dollichon J, Meier F, Meinecke S, Braun M. pandapower—an open-source python tool for convenient modeling, analysis, and optimization of electric power systems. *IEEE Transactions on Power Systems*, 2018, 33(6): 6510–6521
- [55] Donon A, Donnot B, Guyon I, Marot A. Powergraph: A power grid benchmark dataset for graph neural networks. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2024, 1–12
- [56] Li H, Weng Y, Vittal V. Distribution grid topology and parameter estimation using deep-shallow neural network with physical consistency. *IEEE Transactions on Smart Grid*, 2024, 15(1): 1–12
- [57] W3C. Resource description framework (rdf): Concepts and abstract syntax. *W3C Recommendation*, 2004
- [58] Liu F T, Kai M T, Zhou Z H. Isolation forest. *Proceedings of the IEEE International Conference on Data Mining (ICDM)*, 2008, 0: 413–422
- [59] Rousseeuw P J, Van Driessen K. A fast algorithm for the minimum covariance determinant estimator. *Technometrics*, 1999, 41(3): 212–223
- [60] Karimi H S, Natarajan B. Joint topology identification and state estimation in unobservable distribution grids. *IEEE Transactions on Smart Grid*, 2021, 12(6): 4954–4964
- [61] Donon A, Huang B, Lelarge M, Ernst D. Graph neural solver of power systems. In: *IEEE International Conference on Smart Grid Technologies (SmartGridComm)*. 2019
- [62] Li X, Li X, Liu Q, Jiang L, Jia S, Hu R. Digital twin-driven graph attention network for fault diagnosis in power transformer. *Nature Communications*, 2026, 17(1): 3673
- [63] Ministry of Housing and Urban-Rural Development of China. Standard for design of cable lines of electric power engineering, 2018. GB/T 50217-2018, Chinese National Standard
- [64] Standardization Administration of China. Power quality—voltage deviation in power supply systems, 2008. GB/T 12325-2008, Chinese National Standard

- [65] Wang H, Liu Y, Zhang M. Spatiotemporal graph convolutional neural network forecasting-aided state estimation. *IEEE Internet of Things Journal*, 2024, 11(8): 14234–14248
- [66] Donon A, Huang B, Lelarge M, Ernst D. Distributed power system state estimation using graph convolutional neural networks. *IEEE Transactions on Power Systems*, 2021, 36(5): 1–10
- [67] Debelle T, Sohrab F, Barnaghi P. Anomaly detection in smart power grids with graph-regularized MS-SVDD: A multimodal subspace learning approach. *Scientific Reports*, 2026, 16: 1–15
- [68] Uddin M J, Olojede D R, Jacob R A, Coskunuzer B, Zhang J. MP-Grid: Detecting power grid outages with topological machine learning. *Applied Energy*, 2026, 410: 126152
- [69] IEEE . IEEE Standard for Electric Power Systems Communications-Distributed Network Protocol (DNP3), 2012
- [70] International Electrotechnical Commission . Communication networks and systems for power utility automation, 2013. International Standard